

Data Science Project - Predicting Fraud in Insurance Claims Using Machine Learning

1. Aims, objectives and plan

1.1 Aims and Objectives of the Project

The primary aim of this project is to develop a unbiased predictive model to assist an insurance company in identifying potential fraudulent claims while minimising customer loss due to false referrals. The model should meet the following key business and technical objectives:

1. **Business Perspective:**

- Using a model in which the gross profit, amounts to double the claims, and their customer base has 10% claim rate, find how many customers they have.
- Given the model, what the charge for each policy.
- Quantify the financial impact of prediction errors using a pricing model.

2. **Technical Perspective:**

- Preprocess the data removing any noisy attributes, synonymous attributes, feature selection/extraction, collinearity, rescaling, missing and duplicate values and class imbalance
- Build a robust data analysis pipeline using Python and Jupyter Notebook.
- Apply two advanced machine learning techniques to solve the binary classification problem.
- Apply cross validation to the techniques and performance metrics to compare results
- Test the techniques on the test data and evaluate the performance of the models using appropriate metrics.
- Ensure balanced performance with an error rate of less than 5%.

```
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
from datetime import datetime, timedelta

# Define tasks, start dates, and durations
tasks = [
    "Understand Case Study and Objectives",
    "Merge and Preprocess Data",
    "Feature Engineering",
    "Model 1 Development (e.g., Logistic Regression)",
    "Model 2 Development (e.g., Random Forest)",
    "Model Evaluation and Comparison",
    "Financial Impact Analysis",
    "Report Writing",
```

```

    "Explanatory Video Creation",
    "Final Review and Submission"
]

start_dates = [
    datetime(2024, 12, 21),
    datetime(2024, 12, 23),
    datetime(2024, 12, 25),
    datetime(2024, 12, 28),
    datetime(2024, 12, 30),
    datetime(2025, 1, 2),
    datetime(2025, 1, 3),
    datetime(2025, 1, 4),
    datetime(2025, 1, 5),
    datetime(2025, 1, 6)
]

durations = [2, 3, 3, 3, 3, 2, 1, 2, 1, 1] # Durations in days

# Calculate end dates
end_dates = [start_dates[i] + timedelta(days=durations[i]) for i in
range(len(tasks))]

#DataFrame for plotting
gantt_data = pd.DataFrame({
    "Task": tasks,
    "Start": start_dates,
    "End": end_dates
})

# Plotting the Gantt chart
fig, ax = plt.subplots(figsize=(10, 6))

# Iterate through tasks to plot
for i, task in enumerate(gantt_data["Task"]):
    ax.barh(task, (gantt_data["End"][i] - gantt_data["Start"]
[i]).days, left=gantt_data["Start"][i], color="skyblue")

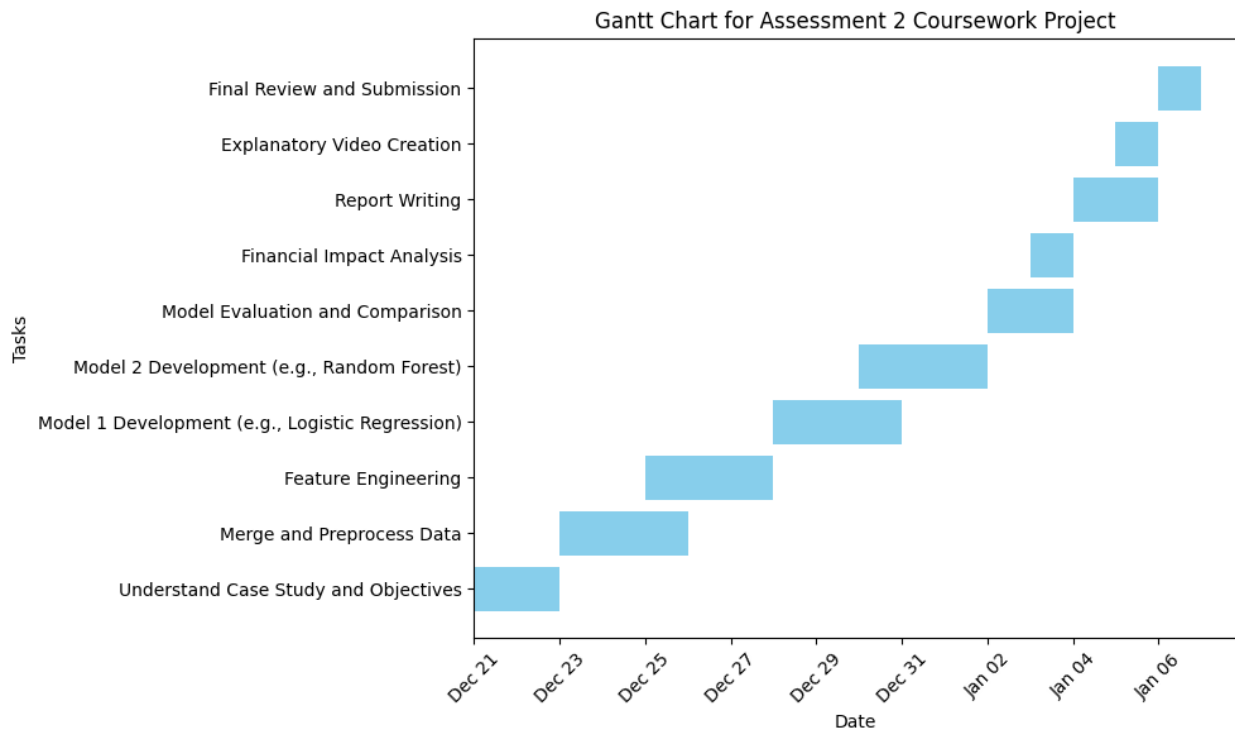
# Format the x-axis for dates
ax.xaxis.set_major_formatter(mdates.DateFormatter("%b %d"))
ax.xaxis.set_major_locator(mdates.DayLocator(interval=2))
plt.xticks(rotation=45)

# labels and title
ax.set_xlabel("Date")
ax.set_ylabel("Tasks")
ax.set_title("Gantt Chart for Assessment 2 Coursework Project")

# Display the chart

```

```
plt.tight_layout()
plt.show()
```



2. Understanding the case study

1.1 Key Points Identified in the Case and Plan to Address Them

1. Objective of the Model

- The primary goal is to develop a predictive model that can accurately identify fraudulent insurance claims while minimising financial losses due to false positives and false negatives.
- Plan: Implement a binary classification model that can distinguish between genuine and fraudulent claims with high precision and recall. Evaluate two machine learning techniques; a tree based model such as gradientboost and multi-layer perceptron neural network technique. Use balanced performance metrics to evaluate each of the model's effectiveness in minimising errors.

1. Dataset Quality and Preprocessing

- The dataset may contain missing values, noisy attributes, redundant information, and class imbalance, which can affect model performance.
- Plan: Perform thorough data preprocessing to handle missing values, remove noisy and redundant attributes, address class imbalance, and rescale numerical features as necessary. Implement feature selection/extraction to reduce dimensionality and collinearity.

1. Impact of False Positives and False Negatives

- False positives can lead to genuine claims being flagged as fraudulent, resulting in customer dissatisfaction and financial loss. False negatives allow fraudulent claims to pass undetected, leading to direct financial losses.
 - Plan: Implement cost-sensitive learning to incorporate the financial impact of errors into model training. Use balanced performance metrics such as the F1-score, precision, and recall to evaluate models.
1. **Model Validation and Evaluation**
 - The model should be validated using cross-validation techniques to ensure robustness and generalisation to unseen data. Performance metrics such as accuracy, precision, recall, F1-score, and ROC-AUC should be used to evaluate model effectiveness.
 - Plan: Apply k-fold cross-validation to evaluate model performance on training data. Use appropriate performance metrics to compare the effectiveness of the two machine learning techniques. Test the final models on the test data to assess their performance in a real-world scenario.

3. Pre-processing applied

3.1 Merging, Pivoting and melting if necessary

```
import pandas as pd
import numpy as np

claims = pd.read_csv(r'archive 2\archive\TrainData\TrainData\
Train_Claim.csv')
demographics = pd.read_csv(r'archive 2\archive\TrainData\TrainData\
Train_Demographics.csv')
policy = pd.read_csv(r'archive 2\archive\TrainData\TrainData\
Train_Policy.csv')
vehicle = pd.read_csv(r'archive 2\archive\TrainData\TrainData\
Train_Vehicle.csv')
target = pd.read_csv(r'archive 2\archive\TrainData\TrainData\
Traindata_with_Target.csv')
#withouttargets = pd.read_csv(r'archive 2\archive\TrainData\TrainData\
Traindata_withoutTarget.csv')

df_merged = (
    claims
    .merge(demographics, on="CustomerID", how="left")
    .merge(policy, on="CustomerID", how="left")
    .merge(vehicle, on="CustomerID", how="left")
    .merge(target, on="CustomerID", how="left")
)

print(df_merged.shape)
df_merged.head()

(115344, 40)
```

	CustomerID	DateOfIncident	TypeOfIncident	TypeOfCollission
\				
0	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
1	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
2	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
3	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
4	Cust10001	2015-02-02	Multi-vehicle Collision	Side Collision
	SeverityOfIncident	AuthoritiesContacted	IncidentState	IncidentCity
\				
0	Total Loss	Police	State7	City1
1	Total Loss	Police	State7	City1
2	Total Loss	Police	State7	City1
3	Total Loss	Police	State7	City1
4	Total Loss	Police	State7	City5
	IncidentAddress	IncidentTime	...	DateOfPolicyCoverage
\				
0	Location 1311	17	...	1998-10-25
1	Location 1311	17	...	1998-10-25
2	Location 1311	17	...	1998-10-25
3	Location 1311	17	...	1998-10-25
4	Location 1311	10	...	2000-11-15
	InsurancePolicyState	Policy_CombinedSingleLimit	Policy_Deductible	
\				
0	State1	100/300		1000
1	State1	100/300		1000
2	State1	100/300		1000
3	State1	100/300		1000
4	State1	100/300		1000
	PolicyAnnualPremium	UmbrellaLimit	InsuredRelationship	VehicleAttribute
\				
0	1632.73	0	not-in-family	VehicleID
1	1632.73	0	not-in-family	

VehicleModel			
2	1632.73	0	not-in-family
VehicleYOM			
3	1632.73	0	not-in-family
VehicleMake			
4	1255.19	0	not-in-family
VehicleYOM			

	VehicleAttributeDetails	ReportedFraud
0	Vehicle26917	N
1	A5	N
2	2008	N
3	Audi	N
4	2006	N

[5 rows x 40 columns]

`df_merged.head()` # Verifying column names and keys to merge datasets and identify potential issues with nomenclature

	CustomerID	DateOfIncident	TypeOfIncident	TypeOfCollission
\				
0	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
1	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
2	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
3	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
4	Cust10001	2015-02-02	Multi-vehicle Collision	Side Collision

	SeverityOfIncident	AuthoritiesContacted	IncidentState
IncidentCity \			
0	Total Loss	Police	State7 City1
1	Total Loss	Police	State7 City1
2	Total Loss	Police	State7 City1
3	Total Loss	Police	State7 City1
4	Total Loss	Police	State7 City5

	IncidentAddress	IncidentTime	...	DateOfPolicyCoverage
\				
0	Location 1311	17	...	1998-10-25
1	Location 1311	17	...	1998-10-25
2	Location 1311	17	...	1998-10-25
3	Location 1311	17	...	1998-10-25

4	Location	1311	10	...	2000-11-15
---	----------	------	----	-----	------------

	InsurancePolicyState	Policy_CombinedSingleLimit	
Policy_Deductible \			
0	State1	100/300	1000
1	State1	100/300	1000
2	State1	100/300	1000
3	State1	100/300	1000
4	State1	100/300	1000

	PolicyAnnualPremium	UmbrellaLimit	InsuredRelationship
VehicleAttribute \			
0	1632.73	0	not-in-family
VehicleID			
1	1632.73	0	not-in-family
VehicleModel			
2	1632.73	0	not-in-family
VehicleYOM			
3	1632.73	0	not-in-family
VehicleMake			
4	1255.19	0	not-in-family
VehicleYOM			

	VehicleAttributeDetails	ReportedFraud
0	Vehicle26917	N
1	A5	N
2	2008	N
3	Audi	N
4	2006	N

[5 rows x 40 columns]

*#I will create a key for each instance in the dataset by concatenating 'CustomerID' and 'DateOfIncident' to uniquely identify each claim.
This will help in pivoting/melting dataset and identifying potential issues with duplicates*

```
df_merged['CustomerID'] = df_merged['CustomerID'].astype(str)
df_merged['DateOfIncident'] =
pd.to_datetime(df_merged['DateOfIncident'])
df_merged['Unique_ID'] = df_merged['CustomerID'] +
df_merged['DateOfIncident'].dt.strftime('%Y%m%d')
df_merged.head()
```

CustomerID	DateOfIncident	TypeOfIncident	TypeOfCollission
\			

0	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
1	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
2	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
3	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
4	Cust10001	2015-02-02	Multi-vehicle Collision	Side Collision

SeverityOfIncident AuthoritiesContacted IncidentState IncidentCity \				
0	Total Loss	Police	State7	City1
1	Total Loss	Police	State7	City1
2	Total Loss	Police	State7	City1
3	Total Loss	Police	State7	City1
4	Total Loss	Police	State7	City5

IncidentAddress IncidentTime ... InsurancePolicyState \				
0	Location 1311	17	...	State1
1	Location 1311	17	...	State1
2	Location 1311	17	...	State1
3	Location 1311	17	...	State1
4	Location 1311	10	...	State1

Policy_CombinedSingleLimit Policy_Deductible PolicyAnnualPremium \			
0	100/300	1000	1632.73
1	100/300	1000	1632.73
2	100/300	1000	1632.73
3	100/300	1000	1632.73
4	100/300	1000	1255.19

UmbrellaLimit InsuredRelationship VehicleAttribute \		
0	0	not-in-family VehicleID
1	0	not-in-family VehicleModel
2	0	not-in-family VehicleYOM
3	0	not-in-family VehicleMake
4	0	not-in-family VehicleYOM

VehicleAttributeDetails ReportedFraud		Unique_ID
0	Vehicle26917	N Cust1000020150203
1	A5	N Cust1000020150203
2	2008	N Cust1000020150203
3	Audi	N Cust1000020150203
4	2006	N Cust1000120150202


```
[5 rows x 41 columns]
```

```
#create table of distinct values in each column for categorical columns
```

```
distinct_values_cat =  
df_merged.select_dtypes(include='object').nunique()  
distinct_values_cat
```

CustomerID	28836
TypeOfIncident	4
TypeOfCollission	4
SeverityOfIncident	4
AuthoritiesContacted	4
IncidentState	7
IncidentCity	7
IncidentAddress	1000
PropertyDamage	3
Witnesses	5
PoliceReport	3
AmountOfTotalClaim	21976
InsuredGender	2
InsuredEducationLevel	7
InsuredOccupation	14
InsuredHobbies	20
Country	1
DateOfPolicyCoverage	6779
InsurancePolicyState	3
Policy_CombinedSingleLimit	9
InsuredRelationship	6
VehicleAttribute	4
VehicleAttributeDetails	28911
ReportedFraud	2
Unique_ID	28836

```
dtype: int64
```

```
#print the first five values of the columns 'VehicleAttribute' and 'VehicleAttributeDetails'
```

```
vehicle_attributes = df_merged[['VehicleAttribute',  
'VehicleAttributeDetails']]  
vehicle_attributes.head()
```

	VehicleAttribute	VehicleAttributeDetails
0	VehicleID	Vehicle26917
1	VehicleModel	A5
2	VehicleYOM	2008
3	VehicleMake	Audi
4	VehicleYOM	2006

The data is “long” rather than “wide.” Instead of having one row per vehicle with columns like VehicleID, VehicleModel, VehicleYOM, VehicleMake, we see multiple rows for a single vehicle record, each row describing one attribute. I'll pivot the data to make it “wide” so that each vehicle record is represented by a single row with all attributes as columns.

#I will write key processes as methods to be reused in the pipeline if needed

```
def group_vehicle_attributes_by_customer_and_date(df):
    """
    Groups rows by (CustomerID, DateOfIncident) and includes:
    - Aggregated leftover columns (e.g., TypeOfIncident,
    SeverityOfIncident, etc.)
    - Pivoted VehicleAttribute columns (VehicleID, VehicleModel,
    VehicleYOM, VehicleMake, etc.)

    Returns a single DataFrame with all columns plus pivoted vehicle
    attributes.
    """

    #Identify columns that belong to VehicleAttribute pivoting
    # and the columns we want to aggregate by (CustomerID,
    DateOfIncident, etc.).
    pivot_cols = ["VehicleAttribute", "VehicleAttributeDetails"]
    group_cols = ["CustomerID", "DateOfIncident"]

    # The leftover columns are the ones we want to keep as-is.

    leftover_cols = [
        c for c in df.columns
        if c not in (group_cols + pivot_cols)
    ]

    # Aggregate leftover columns by taking the first row in each
    (CustomerID, DateOfIncident) group

    df_agg = (
        df.groupby(group_cols, as_index=False)[leftover_cols]
        .agg("first")
    )

    # Pivot the vehicle attributes so each attribute becomes a column

    pivoted = (
        df.pivot_table(
            index=group_cols,
            columns="VehicleAttribute",
            values="VehicleAttributeDetails",
            aggfunc=lambda x: ",".join(x) # combine multiple values
            with commas
        )
    )
```

```

    )
    .reset_index()
)

#flatten pivot table if needed
pivoted.columns.name = None # remove top-level name
pivoted = pivoted.rename_axis(None, axis=1) # remove index name
if needed

# Merge the aggregated leftover columns with the pivoted vehicle
attributes
final_df = pd.merge(df_agg, pivoted, on=group_cols, how="left")

# Reorder columns so that group_cols come first, then
leftover_cols, then pivoted columns
pivoted_attribs = sorted(set(pivoted.columns) - set(group_cols))
final_cols = group_cols + leftover_cols + pivoted_attribs
final_cols = [c for c in final_cols if c in final_df.columns] #
only keep existing
final_df = final_df[final_cols]

return final_df

df_pivot = group_vehicle_attributes_by_customer_and_date(df_merged)
df_pivot

```

	CustomerID	DateOfIncident	TypeOfIncident	
TypeOfCollision \				
0	Cust10000	2015-02-03	Multi-vehicle Collision	Side Collision
1	Cust10001	2015-02-02	Multi-vehicle Collision	Side Collision
2	Cust10002	2015-01-15	Single Vehicle Collision	Side Collision
3	Cust10003	2015-01-19	Single Vehicle Collision	Side Collision
4	Cust10004	2015-01-09	Single Vehicle Collision	Rear Collision
...
...				
28831	Cust9993	2015-01-24	Vehicle Theft	
?				
28832	Cust9994	2015-02-09	Single Vehicle Collision	Front Collision
28833	Cust9996	2015-01-28	Single Vehicle Collision	Rear Collision
28834	Cust9997	2015-01-28	Single Vehicle Collision	Rear Collision
28835	Cust9999	2015-01-13	Single Vehicle Collision	Front Collision

IncidentCity \	SeverityOfIncident	AuthoritiesContacted	IncidentState
0	Total Loss	Police	State7
City1			
1	Total Loss	Police	State7
City5			
2	Minor Damage	Other	State8
City6			
3	Minor Damage	Other	State9
City6			
4	Minor Damage	Fire	State8
City6			
...
..			
28831	Trivial Damage	Police	State9
City4			
28832	Minor Damage	Fire	State8
City3			
28833	Minor Damage	Fire	State9
City7			
28834	Minor Damage	Ambulance	State9
City3			
28835	Total Loss	Other	State8
City3			

	IncidentAddress	IncidentTime	...	Policy_Deductible \
0	Location 1311	17	...	1000
1	Location 1311	10	...	1000
2	Location 2081	22	...	617
3	Location 2081	22	...	722
4	Location 1695	10	...	500
...
28831	Location 1890	3	...	655
28832	Location 2097	17	...	1089
28833	Location 1452	1	...	787
28834	Location 1876	1	...	780
28835	Location 1874	9	...	2000

PolicyAnnualPremium	UmbrellaLimit	InsuredRelationship
ReportedFraud \		
0	1632.73	0 not-in-family
N		
1	1255.19	0 not-in-family
N		
2	1373.38	0 wife
N		
3	1337.60	0 own-child
N		
4	1353.73	4279863 unmarried

```

N
...
...
28831      1276.01      0      unmarried
N
28832      1273.38      0      unmarried
N
28833      1380.92      3448735      own-child
N
28834      1389.29      3364301      not-in-family
N
28835      928.43      2909175      husband
N
      Unique_ID      VehicleID      VehicleMake      VehicleModel
VehicleYOM
0      Cust1000020150203      Vehicle26917      Audi      A5
2008
1      Cust1000120150202      Vehicle15893      Audi      A5
2006
2      Cust1000220150115      Vehicle5152      Volkswagen      Jetta
1999
3      Cust1000320150119      Vehicle37363      Volkswagen      Jetta
2003
4      Cust1000420150109      Vehicle28633      Toyota      CRV
2010
...
...
28831      Cust999320150124      Vehicle13568      Suburu      Impreza
2007
28832      Cust999420150209      Vehicle19810      Jeep      Wrangler
2003
28833      Cust999620150128      Vehicle3339      Suburu      Legacy
2004
28834      Cust999720150128      Vehicle10240      Suburu      Forrester
2004
28835      Cust999920150113      Vehicle39163      Suburu      E400
2007
[28836 rows x 43 columns]

```

3.2 Identifying and dealing with class imbalance, if necessary

Data Splitting and Validation

To avoid data leakage, I will split the data into training and test sets before performing any preprocessing steps. I will use a stratified split to ensure that the class distribution is maintained in both sets. I will also perform a preliminary check to ensure that the split was successful and that the class distribution is balanced in both sets.

```

# split the data

df = df_pivot
X = df.drop(columns=["ReportedFraud"])
y = df["ReportedFraud"]

from sklearn.model_selection import train_test_split

# Train (80%) vs. Test (20%)
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.20,           # 20% for test
    random_state=42,         # ensures reproducibility
    stratify=y,              # maintain class ratio, e.g., fraud/non-
    fraud
)

print(X_train.shape)
print(y_train.shape)

print(X_test.shape)
print(y_test.shape)

(23068, 42)
(23068,)
(5768, 42)
(5768,)

# percentage of each class in y_train
fraud_train_counts = y_train.value_counts()
fraud_train_percent = y_train.value_counts(normalize=True) * 100

print("Fraud Distribution in the Training Set:")
for fraud_class in fraud_train_counts.index:
    count = fraud_train_counts[fraud_class]
    percent = fraud_train_percent[fraud_class]
    print(f" {fraud_class}: {count} instances ({percent:.2f}%)")

Fraud Distribution in the Training Set:
N: 16840 instances (73.00%)
Y: 6228 instances (27.00%)

# Percentage of each class in y_test
fraud_test_counts = y_test.value_counts()
fraud_test_percent = y_test.value_counts(normalize=True) * 100

print("Fraud Distribution in the Test:")
for fraud_class in fraud_test_counts.index:

```

```
count = fraud_test_counts[fraud_class]
percent = fraud_test_percent[fraud_class]
print(f" {fraud_class}: {count} instances ({percent:.2f}%)" )
```

Fraud Distribution in the Test:

N: 4211 instances (73.01%)

Y: 1557 instances (26.99%)

The split suggests that the class distribution is maintained in both the training and test sets, indicating a successful stratified split. The training set contains 80% of the data, while the test set contains 20%. The class distribution is balanced in both sets, with approximately 27% of claims being fraudulent in each set.

3.3 Appropriate feature extraction, if necessary

Dates and Times

The dates and times can be extracted from the 'DateOfIncident' and 'DateOfClaim' columns to create new features such as 'DayOfWeek', 'Month', 'Year', 'Weekday', 'Weekend', 'TimeOfDay', etc. These features can make it more convenient for machine learning models to learn patterns in the data as well as for calculations.

```
def process_incident_and_policy_dates(df,
                                     incident_col='DateOfIncident',
                                     policy_col='DateOfPolicyCoverage',
                                     incident_year_col='IncidentYear', vehicle_yom_col='VehicleYOM'):
    """
```

Converts the incident and policy coverage dates to datetime, then creates columns for year/month/day of incident and calculates policy age in days.

Parameters:

df : pd.DataFrame

The DataFrame containing date columns.

incident_col : str

Column name with the incident date.

policy_col : str

Column name with the policy coverage start date.

Returns:

pd.DataFrame

The original DataFrame with new columns added:

- IncidentYear
- IncidentMonth
- IncidentDay

```

        - PolicyAgeDays
    """

    # Convert incident date to datetime
    df[incident_col] = pd.to_datetime(df[incident_col],
errors='coerce')

    # Extract components of the incident date
    df['IncidentYear'] = df[incident_col].dt.year
    df['IncidentMonth'] = df[incident_col].dt.month
    df['IncidentDay'] = df[incident_col].dt.day

    # Convert policy coverage date to datetime
    df[policy_col] = pd.to_datetime(df[policy_col], errors='coerce')

    # Calculate the number of days between incident and policy
coverage
    df['PolicyAgeDays'] = (df[incident_col] - df[policy_col]).dt.days

    # Convert both columns to numeric (int)
    df[incident_year_col] = pd.to_numeric(df[incident_year_col],
errors='coerce')
    df[vehicle_yom_col] = pd.to_numeric(df[vehicle_yom_col],
errors='coerce')

    # Calculate age at incident in years
    df['AgeAtIncident'] = df[incident_year_col] - df[vehicle_yom_col]

    return df

#apply the function to the training and test data
X_train_grouped = process_incident_and_policy_dates(X_train)
X_test_grouped = process_incident_and_policy_dates(X_test)

#check for data integrity
print(X_train_grouped['Unique_ID'].nunique)
print(X_train_grouped.shape)
print(X_test_grouped.shape)
print(y_train.shape)
print(y_test.shape)

<bound method IndexOpsMixin.nunique of 20936      Cust3542720150227
19307      Cust3344820150208
2609      Cust1315920150117
15420      Cust2870520150225
21007      Cust3551220150110
...
27962      Cust881320150222
10263      Cust2239220150225
11822      Cust2422820150114

```



```

8667      Cust2048020150216
10825      Cust2303720150103
Name: Unique_ID, Length: 23068, dtype: object>
(23068, 47)
(5768, 47)
(23068,)
(5768,)

```

3.4 Dealing with Missing Values (Imputation/Filtering) without Leakage

```

#missing values in the training set
missing_values = X_train_grouped.isnull().sum()
missing_values = missing_values[missing_values > 0]
missing_values

AuthoritiesContacted      2144
InsuredGender              23
Country                    1
dtype: int64

#other categories of missing values such as 'Unknown' or 'MISSEDATA' or '?'
unknown_values = X_train_grouped.isin(['Unknown', 'MISSEDATA',
'?']).sum()
unknown_values = unknown_values[unknown_values > 0]
unknown_values

TypeOfCollission          4104
PropertyDamage            8375
PoliceReport              7830
AmountOfTotalClaim        40
dtype: int64

#other potential missing data categories such as 'Not Available' or 'N/A'
na_values = X_train_grouped.isin(['Not Available', 'N/A']).sum()
na_values = na_values[na_values > 0]
na_values

Series([], dtype: int64)

```

Unifying Missing Values

```

def unify_missing_values(df, extra_missing=None):
    """
    Converts placeholders like '?' and 'MISSEDATA' (and any extra you
    pass) into np.nan.
    Returns a new DataFrame with unified missing values.
    """

```

Parameters:

df : pd.DataFrame
The DataFrame whose missing placeholders you want to standardize.
extra_missing : list of str, optional
Additional string placeholders that should be converted to np.nan.

Returns:

pd.DataFrame
A copy of the original DataFrame, with all missing placeholders replaced by np.nan.

```
"""
# Default placeholders
placeholders = ['?', 'MISSEDDATA']

# passed additional placeholders
if extra_missing is not None:
    placeholders.extend(extra_missing)

# Replace all placeholders with np.nan

df_clean = df.replace(placeholders, np.nan)

return df_clean

#apply the function to the training and test data
X_train_clean = unify_missing_values(X_train_grouped)
X_test_clean = unify_missing_values(X_test_grouped)

print(X_train_clean.shape)
print(X_test_clean.shape)

(23068, 47)
(5768, 47)
```

Check for Missing Values

```
def check_missing_values(df):
    """
    Checks for missing values in a DataFrame and returns a summary
    DataFrame
    containing both the absolute number of missing values and their
    percentage.
```

Parameters

df : pd.DataFrame
The DataFrame to analyse.

Returns

pd.DataFrame

A DataFrame with two columns:

- "Missing Values" (the count of NaNs in each column)*
- "Percentage Missing (%)" (the percentage of NaNs in that column).*

Only columns with missing values are listed.

```
"""
```

```
# Count total NaNs in each column
```

```
missing = df.isnull().sum()
```

```
# Filter out columns that have no missing values
```

```
missing = missing[missing > 0]
```

```
# Calculate the percentage of missing values
```

```
percentage_missing = (missing / len(df)) * 100
```

```
# Create a summary DataFrame
```

```
missing_values_df = pd.DataFrame({
```

```
    "Missing Values": missing,
```

```
    "Percentage Missing (%)": percentage_missing
```

```
})
```

```
return missing_values_df
```

```
missing_values_train_train = check_missing_values(X_train_clean)
```

```
missing_values_train_train
```

	Missing Values	Percentage Missing (%)
TypeOfCollission	4104	17.790879
AuthoritiesContacted	2144	9.294260
PropertyDamage	8375	36.305705
PoliceReport	7830	33.943125
AmountOfTotalClaim	40	0.173400
InsuredGender	23	0.099705
Country	1	0.004335

```
missing_values_test = check_missing_values(X_test_clean)
```

```
missing_values_test
```

	Missing Values	Percentage Missing (%)
TypeOfCollission	1058	18.342580
AuthoritiesContacted	548	9.500693
PropertyDamage	2084	36.130374
PoliceReport	1975	34.240638

AmountOfTotalClaim	10	0.173370
InsuredGender	7	0.121359
Country	1	0.017337

Imputation Strategy

In order to develop an imputation strategy, I will see if gender has any direct bearing on fields such as hobbies and interests. If there is a correlation, the gender can be inferred from the hobbies and interests.

```
# Group by Hobbies and Gender, count number of rows
hobby_gender_counts = (
    X_train_clean
    .groupby(['InsuredHobbies', 'InsuredGender'])
    .size()
    .unstack(fill_value=0) # Convert to columns = ['F', 'M'], fill
    missing with 0
)

# Convert counts to percentages (row-wise)
hobby_gender_percent =
hobby_gender_counts.div(hobby_gender_counts.sum(axis=1), axis=0) * 100

# Print or inspect
print(hobby_gender_percent)
```

InsuredGender	FEMALE	MALE
InsuredHobbies		
base-jumping	47.147651	52.852349
basketball	46.650426	53.349574
board-games	53.019538	46.980462
bungee-jumping	54.829545	45.170455
camping	52.193309	47.806691
chess	57.113188	42.886812
cross-fit	51.717370	48.282630
dancing	59.756098	40.243902
exercise	61.624204	38.375796
golf	54.402790	45.597210
hiking	51.219512	48.780488
kayaking	65.495706	34.504294
movies	51.915240	48.084760
paintball	57.528090	42.471910
polo	48.805147	51.194853
reading	54.838710	45.161290
skydiving	57.592093	42.407907
sleeping	55.730810	44.269190
video-games	55.460017	44.539983
yachting	46.740995	53.259005

```

# Group by Occupation and Gender, get counts
occ_gender_counts = (
    X_train_clean
    .groupby(['InsuredOccupation', 'InsuredGender'])
    .size()
    .unstack(fill_value=0)
)

# Convert to percentages
occ_gender_percent =
occ_gender_counts.div(occ_gender_counts.sum(axis=1), axis=0) * 100

# Print or inspect
print(occ_gender_percent)

```

InsuredGender	FEMALE	MALE
InsuredOccupation		
adm-clerical	58.942559	41.057441
armed-forces	50.945699	49.054301
craft-repair	54.821742	45.178258
exec-managerial	48.711944	51.288056
farming-fishing	58.550039	41.449961
handlers-cleaners	52.468619	47.531381
machine-op-inspct	51.351351	48.648649
other-service	52.918538	47.081462
priv-house-serv	51.875000	48.125000
prof-specialty	56.825397	43.174603
protective-serv	56.350365	43.649635
sales	57.017544	42.982456
tech-support	58.810198	41.189802
transport-moving	52.797619	47.202381

No evident correlation for gender with hobbies and interests. Therefore, I will use the mode imputation strategy to fill missing values

Observations:

Categorical and numerical features with low or medium percentage missing values can be need not be dropped and can be imputed with a separate category called 'Unknown' or the mode value.

Categorical Values: AmountOfTotalClaim, InsuredGender, Country have very low missing rates (<1%) TypeOfCollission has moderate missing (~18%) AuthoritiesContacted has moderate missing (~9%)

Numerical Amount of Total Claim has very low missing rates (<1%) and can be imputed with the median value.

Categories with high missing values: PropertyDamage and PoliceReport each have over 30% missing. I can evaluate if losing an feature is worse than risking biased imputations. In prediction

fraud, these features may be important, so we will impute them with a separate category called 'Unknown', as the missing values may carry information and are not above 40%

```
#change AmountOfTotalClaim to float
df = X_train_clean
df['AmountOfTotalClaim'] = df['AmountOfTotalClaim'].astype(float)

df = X_test_clean
df['AmountOfTotalClaim'] = df['AmountOfTotalClaim'].astype(float)
```

3.5 Dealing with duplicate values, if necessary

```
# Based on unique ID, there are no duplicate rows in the dataset.
```

3.6 Categorical and Numerical Encoding if Necessary

I will identify categorical and numerical columns in the dataset and encode them appropriately. Categorical columns will be one-hot encoded, while numerical columns can be scaled if needed. I will also check for any remaining missing values and handle them accordingly. I can also use frequency encoding for categorical columns with high cardinality to reduce dimensionality.

```
def get_numeric_and_categorical_columns(df):
    """
    Identifies numeric and categorical columns in the DataFrame.

    Parameters
    df : pd.DataFrame
        The DataFrame to analyze.

    Returns
    numerical_columns : list
        A list of column names that have int or float data types.
    categorical_columns : list
        A list of column names that have object data types.

    """
    numerical_columns = df.select_dtypes(include=[int,
float]).columns.tolist()
    categorical_columns =
df.select_dtypes(include=['object']).columns.tolist()

    return numerical_columns, categorical_columns

num_cols, cat_cols =
get_numeric_and_categorical_columns(X_train_clean)
print("Numeric columns:", num_cols)
print("Categorical columns:", cat_cols)
```

```
Numeric columns: ['IncidentTime', 'NumberOfVehicles',  
'BodilyInjuries', 'AmountOfTotalClaim', 'AmountOfInjuryClaim',  
'AmountOfPropertyClaim', 'AmountOfVehicleDamage', 'InsuredAge',  
'InsuredZipCode', 'CapitalGains', 'CapitalLoss',  
'InsurancePolicyNumber', 'CustomerLoyaltyPeriod', 'Policy_Deductible',  
'PolicyAnnualPremium', 'UmbrellaLimit', 'VehicleYOM', 'IncidentYear',  
'IncidentMonth', 'IncidentDay', 'PolicyAgeDays', 'AgeAtIncident']  
Categorical columns: ['CustomerID', 'TypeOfIncident',  
'TypeOfCollission', 'SeverityOfIncident', 'AuthoritiesContacted',  
'IncidentState', 'IncidentCity', 'IncidentAddress', 'PropertyDamage',  
'Witnesses', 'PoliceReport', 'InsuredGender', 'InsuredEducationLevel',  
'InsuredOccupation', 'InsuredHobbies', 'Country',  
'InsurancePolicyState', 'Policy_CombinedSingleLimit',  
'InsuredRelationship', 'Unique_ID', 'VehicleID', 'VehicleMake',  
'VehicleModel']
```

```
from sklearn.impute import SimpleImputer  
  
numeric_imputer = SimpleImputer(strategy='mean')  
  
categorical_imputer = SimpleImputer(strategy='most_frequent')  
  
from sklearn.compose import ColumnTransformer  
from sklearn.pipeline import Pipeline  
  
# Define separate pipelines for numeric and categorical  
numeric_pipeline = Pipeline([  
    ('imputer', SimpleImputer(strategy='mean'))  
    # Add other numeric transforms if needed, e.g. StandardScaler  
)  
  
categorical_pipeline = Pipeline([  
    ('imputer', SimpleImputer(strategy='most_frequent'))  
    # Add other categorical transforms if needed, e.g. OneHotEncoder  
)  
  
# Combine them in a ColumnTransformer  
preprocessor = ColumnTransformer([  
    ('num', numeric_pipeline, num_cols),  
    ('cat', categorical_pipeline, cat_cols)  
)  
  
  
# Fit on training data  
preprocessor.fit(X_train_clean)  
  
# Transform training data  
X_train_imputed = preprocessor.transform(X_train_clean)  
  
# Transform test data
```

```
X_test_imputed = preprocessor.transform(X_test_clean)

#make dataframe
X_train_imputed = pd.DataFrame(X_train_imputed, columns=num_cols +
cat_cols)
X_test_imputed = pd.DataFrame(X_test_imputed, columns=num_cols +
cat_cols)

check_missing_values(X_train_imputed)
```

	Missing Values	Percentage Missing (%)
AuthoritiesContacted	2144	9.294260
InsuredGender	23	0.099705
Country	1	0.004335

```
#add 'missing' category to categorical columns with missing values
X_train_imputed.fillna('missing', inplace=True)
X_test_imputed.fillna('missing', inplace=True)
```

```
C:\Users\saqib\AppData\Local\Temp\ipykernel_30288\1039554909.py:2:
FutureWarning: Downcasting object dtype arrays
on .fillna, .ffill, .bfill is deprecated and will change in a future
version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
X_train_imputed.fillna('missing', inplace=True)
C:\Users\saqib\AppData\Local\Temp\ipykernel_30288\1039554909.py:3:
FutureWarning: Downcasting object dtype arrays
on .fillna, .ffill, .bfill is deprecated and will change in a future
version. Call result.infer_objects(copy=False) instead. To opt-in to
the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
X_test_imputed.fillna('missing', inplace=True)
```

```
#check missing values
print("Missing values in the training set after adding 'missing'
category:")
check_missing_values(X_train_imputed)
check_missing_values(X_test_imputed)
```

Missing values in the training set after adding 'missing' category:

```
Empty DataFrame
Columns: [Missing Values, Percentage Missing (%)]
Index: []
```

Frequency Encoding

```
#distinct object columns
```



```

distinct_values =
X_train_clean.select_dtypes(include='object').nunique()
#distinct_values = X_train_clean.nunique()
distinct_values

```

```

CustomerID                23068
TypeOfIncident             4
TypeOfCollission          3
SeverityOfIncident        4
AuthoritiesContacted      4
IncidentState             7
IncidentCity              7
IncidentAddress          1000
PropertyDamage            2
Witnesses                 5
PoliceReport              2
InsuredGender             2
InsuredEducationLevel     7
InsuredOccupation         14
InsuredHobbies            20
Country                   1
InsurancePolicyState      3
Policy_CombinedSingleLimit 9
InsuredRelationship       6
Unique_ID                23068
VehicleID                 23068
VehicleMake               15
VehicleModel              39
dtype: int64

```

```

def frequency_encode_columns(df, columns, drop_original=False,
suffix="_freq"):
    """

```

*Applies frequency encoding to the given columns in a DataFrame.
Creates new columns named <col> + <suffix> by default.*

Parameters

df : pd.DataFrame

The DataFrame with columns to be encoded.

columns : list of str

A list of column names in df to frequency-encode.

drop_original : bool, optional

If True, drop the original columns after encoding. Default is False.

suffix : str, optional

Suffix to append to the new encoded column names. Default is "_freq".

Returns

pd.DataFrame

A DataFrame with the new frequency-encoded columns added.

If drop_original=True, the original columns are removed.

"""

```
df_encoded = df.copy() # Work on a copy
```

```
for col in columns:
```

```
    # 1) Compute frequency counts
```

```
    freq_map = df_encoded[col].value_counts().to_dict()
```

```
    # 2) Create the frequency-encoded column
```

```
    new_col_name = col + suffix
```

```
    df_encoded[new_col_name] = df_encoded[col].map(freq_map)
```

```
    # 3) Optionally drop the original column
```

```
    if drop_original:
```

```
        df_encoded.drop(columns=[col], inplace=True)
```

```
return df_encoded
```

```
#apply frequency encoding to the 'IncidentAddress' column
```

```
X_train_fre_encoded = frequency_encode_columns(X_train_clean,  
["IncidentAddress"], drop_original=True)
```

```
X_test_fre_encoded = frequency_encode_columns(X_test_clean,  
["IncidentAddress"], drop_original=True)
```

3.7 Dealing with correlation and collinearity, if necessary

Correlation and collinearity will be assessed to identify redundant features that may affect model performance.

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
corr_matrix = X_train_fre_encoded.corr(numeric_only=True)
```

```
plt.figure(figsize=(10, 8))
```

```
# Plot the correlation matrix
```

```
im = plt.imshow(corr_matrix, cmap="coolwarm", origin="lower",  
aspect="auto")
```

```
# Add a colorbar
```

```
cbar = plt.colorbar(im)
```

```

cbar.set_label("Correlation Coefficient")

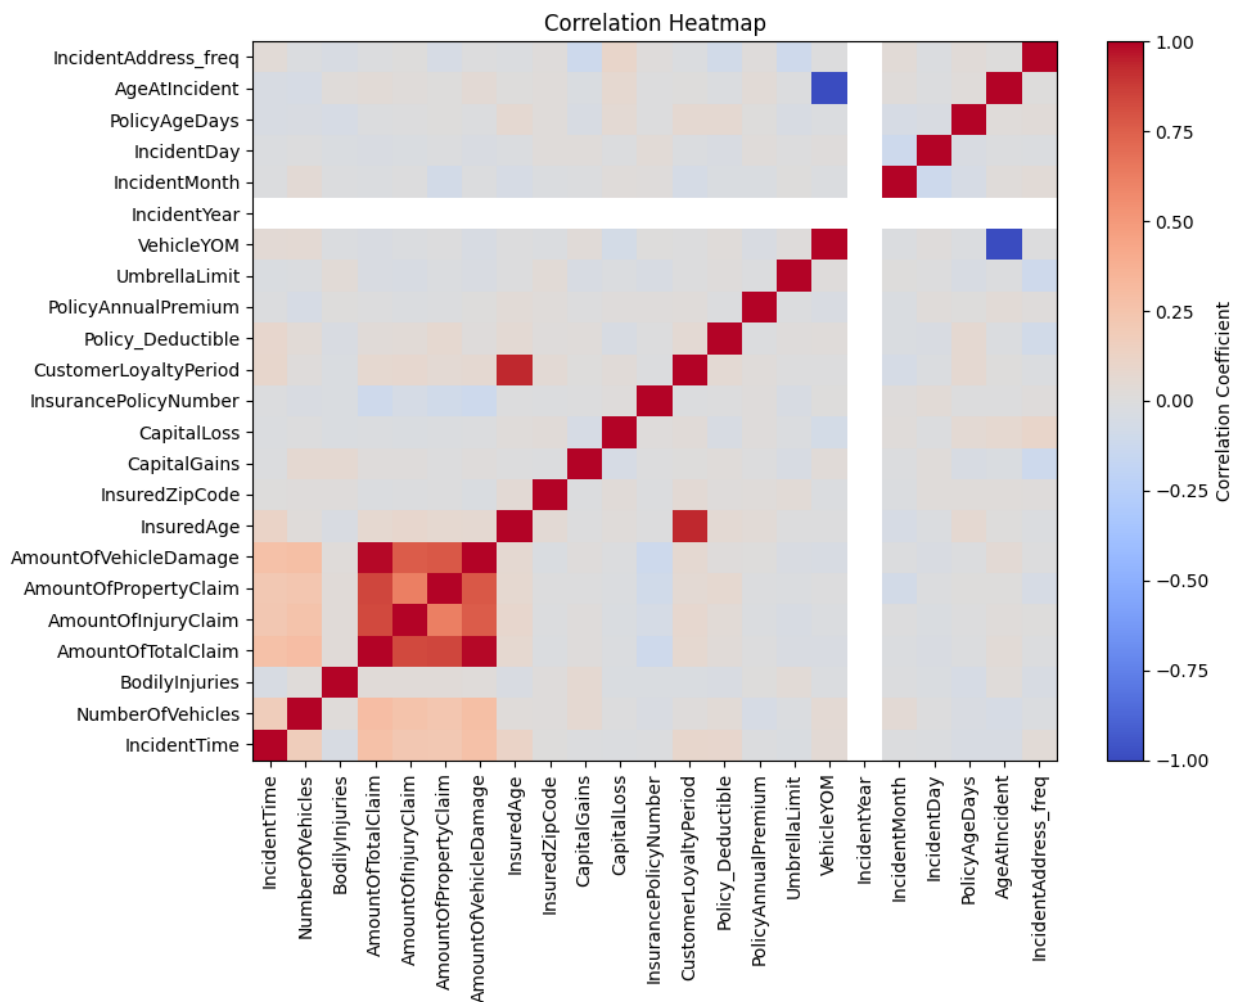
# Add title
plt.title("Correlation Heatmap")

# Set ticks and labels for x and y axes
num_features = len(corr_matrix.columns)
plt.xticks(np.arange(num_features), corr_matrix.columns, rotation=90)
plt.yticks(np.arange(num_features), corr_matrix.columns)

plt.tight_layout()

plt.show()

```



```

# Select numeric columns for correlation matrix calculation
numeric_data = X_train_fre_encoded.select_dtypes(include=[np.number])

```

```

# Calculate correlation matrix of numeric features
corr_matrix = numeric_data.corr().abs()

# Upper triangle of the correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape),
k=1).astype(bool))

# Identify columns with correlation above a chosen threshold
threshold = 0.9
redundant_features = [
    column for column in upper.columns if any(upper[column] >
threshold)
]

print("Highly correlated features:", redundant_features)

```

```

Highly correlated features: ['AmountOfVehicleDamage',
'CustomerLoyaltyPeriod', 'AgeAtIncident']

```

*#following columns can be removed as they have little to no impact on the target e.g. VehicleID,
#or are redundant e.g. Unique_ID
or are highly correlated and represent the same information e.g. Customer loyalty period and age at incident*

#Incident Year has a variance of 0, so it can be removed as well

```

columns_to_remove =
['IncidentYear', 'CustomerID', 'DateOfPolicyCoverage', 'VehicleID', 'Count
ry', 'VehicleYOM', 'CustomerLoyaltyPeriod', 'Unique_ID', 'AmountOfTotalCla
im']

```

#remove columns_to_remove from training and test sets

```

X_train_fre_encoded =
X_train_fre_encoded.drop(columns=columns_to_remove)
X_test_fre_encoded =
X_test_fre_encoded.drop(columns=columns_to_remove)

```

```

import matplotlib.pyplot as plt
import numpy as np

```

```

# Compute correlation matrix
corr_matrix = X_train_fre_encoded.corr(numeric_only=True)

```

```

plt.figure(figsize=(10, 8))

```

```

# Display the correlation matrix using imshow
im = plt.imshow(
    corr_matrix,

```

```

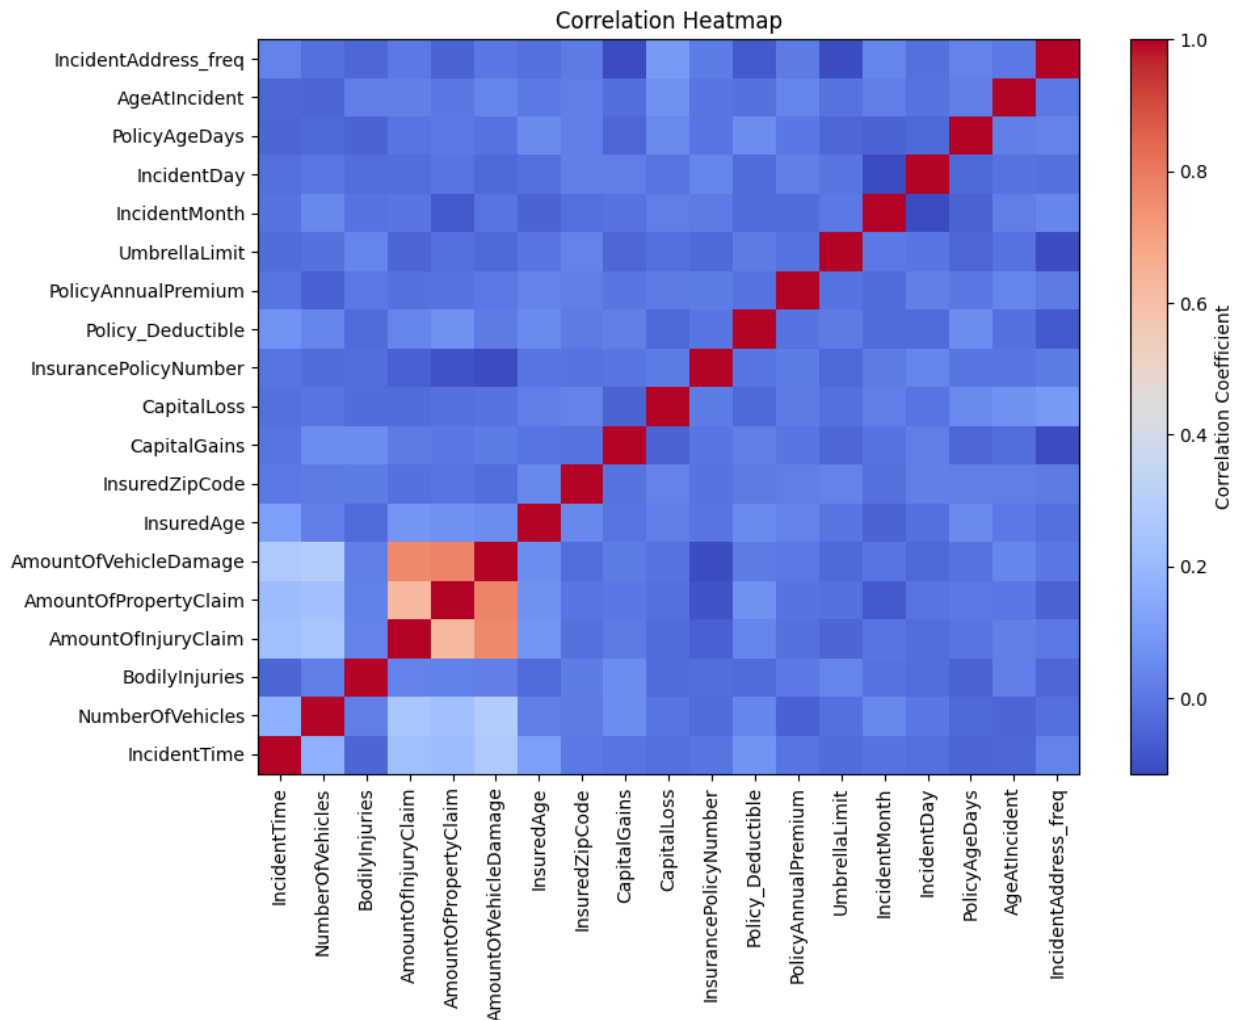
cmap="coolwarm", # same color scheme as in Seaborn
aspect="auto",   # adjusts aspect ratio
origin="lower"   # so that the [0,0] index is at the bottom-left
)

# Add a colorbar to show correlation scale
cbar = plt.colorbar(im)
cbar.set_label("Correlation Coefficient")

# Label the axes with feature names
num_features = len(corr_matrix.columns)
plt.xticks(np.arange(num_features), corr_matrix.columns, rotation=90)
plt.yticks(np.arange(num_features), corr_matrix.columns)

plt.title("Correlation Heatmap")
plt.tight_layout() # avoids label cut-off
plt.show()

```



3.8 Other Pre-processing

Certain fields have a high level of distinct values, and are not suitable for one-hot encoding and can lead to the curse of dimensionality. I will evaluate the cardinality of these fields and decide on an appropriate encoding strategy.

Feature Engineering: Zip Code Encoding

```
df = X_train_fre_encoded

unique_zip_codes = df['InsuredZipCode'].unique().size

print(f"Unique zip codes: {unique_zip_codes}")

Unique zip codes: 995
```

It seems there are 995 unique values for the insured zip code, which is too many for one-hot encoding. This can be dropped or grouped into regions. As region can play an important part in fraud detection, it will be grouped into regions. The codes themselves do correlate with the Indian postal codes as the 4 and 6 digit at the start match the Indian postal code system. There are no codes started with any other digit. The codes will be grouped into regions based on the first two digits of the postal code. I will use Indian region names as labels for the grouping knowing that the codes are there simply for purposes of prediction and not for actual location.

Every code in the list falls into either Maharashtra (40–44), Madhya Pradesh (45–48), or Tamil Nadu (60–66).

```
def pin_to_region(pin_code):
    # Convert to string just in case
    s = str(pin_code)
    first_two = int(s[:2])

    if 40 <= first_two <= 44:
        return "Maharashtra"
    elif 45 <= first_two <= 48:
        return "Madhya Pradesh"
    elif 60 <= first_two <= 66:
        return "Tamil Nadu"
    else:
        return "Other" # or handle other states as needed

#Apply to training set

X_train_fre_encoded['Region'] =
X_train_fre_encoded['InsuredZipCode'].apply(pin_to_region)
X_train_fre_encoded =
X_train_fre_encoded.drop(columns=['InsuredZipCode'])
X_train_fre_encoded['Region'].value_counts()
```

```
Region
Madhya Pradesh    9902
Tamil Nadu       6978
Maharashtra      6188
Name: count, dtype: int64
```

#Apply to test set

```
X_test_fre_encoded['Region'] =
X_test_fre_encoded['InsuredZipCode'].apply(pin_to_region)
X_test_fre_encoded =
X_test_fre_encoded.drop(columns=['InsuredZipCode'])
X_test_fre_encoded['Region'].value_counts()
```

```
Region
Madhya Pradesh    2505
Tamil Nadu       1732
Maharashtra      1531
Name: count, dtype: int64
```

3.9 Variance analysis, if necessary

```
variance_per_column = X_train_fre_encoded.var(numeric_only=True) # or
df.std()
```

variance_per_column

```
IncidentTime          3.791831e+01
NumberOfVehicles      9.637541e-01
BodilyInjuries        6.131465e-01
AmountOfInjuryClaim   1.943763e+07
AmountOfPropertyClaim 1.899265e+07
AmountOfVehicleDamage 3.217226e+08
InsuredAge            6.418861e+01
CapitalGains          7.636088e+08
CapitalLoss           7.770723e+08
InsurancePolicyNumber 1.233447e+08
Policy_Deductible     2.981809e+05
PolicyAnnualPremium   4.998944e+04
UmbrellaLimit         3.904602e+12
IncidentMonth         2.660365e-01
IncidentDay           5.781666e+01
PolicyAgeDays         5.717444e+06
AgeAtIncident         2.822925e+01
IncidentAddress_freq  1.049851e+02
dtype: float64
```

- IncidentYear has 0 variance and can be dropped as it does not provide any information to the model.
- InsurancePolicyNumber can also be dropped as it provides no information to the model.
- Small variance columns such as BodilyInjuries and Number of Vehicles Involved will be kept as they may still provide valuable information for the model.

```
#drop InsurancePolicyNumber and IncidentYear
```

```
X_train_filled =
X_train_fre_encoded.drop(columns=['InsurancePolicyNumber', 'DateOfIncident'])
X_test_filled =
X_test_fre_encoded.drop(columns=['InsurancePolicyNumber', 'DateOfIncident'])
```

3.10 Scaling, without leaking, if necessary

This will depend on technique as tree based models for examples do not require scaling. However, scaling can be beneficial for other models such as SVM, k-NN, and neural networks. I will scale the numerical features as needed in the pipeline to ensure that all features are on the same scale.

```
df_num = X_train_filled.select_dtypes(include=[np.number])
```

```
df_num
```

	IncidentTime	NumberOfVehicles	BodilyInjuries
AmountOfInjuryClaim \			
20936	4	3	2
10529			
19307	19	3	2
4978			
2609	8	1	2
347			
15420	4	2	2
14503			
21007	14	1	0
6367			
...
...			
27962	10	3	0
5591			
10263	21	3	2
7590			
11822	16	3	0
12538			
8667	21	3	0
8729			
10825	12	1	2
7589			

CapitalGains \	AmountOfPropertyClaim	AmountOfVehicleDamage	InsuredAge
209360	5901	34135	44
193070	4710	43613	40
26090	693	2882	38
1542045600	14503	59684	33
210070	12735	43169	47
...
279620	10472	41889	39
1026369200	4863	33917	32
1182254600	11212	48826	43
866733500	8729	38353	45
108250	5233	28000	30

UmbrellaLimit \	CapitalLoss	Policy_Deductible	PolicyAnnualPremium
209360	0	1000	1072.38
193070	-40600	1000	1441.49
26090	-43400	1507	1046.00
154203026658	-61400	500	1216.26
210070	-55600	1181	1493.96
...
279620	-71700	500	1349.37
102630	-36900	2000	1438.39
118220	-45500	638	851.46
86670	0	663	1298.81
10825	0	500	1282.72

0

	IncidentMonth	IncidentDay	PolicyAgeDays	AgeAtIncident	\
20936	2	27	4526	7	
19307	2	8	8092	9	
2609	1	17	3922	14	
15420	2	25	6974	11	
21007	1	10	8370	10	
...	
27962	2	22	466	4	
10263	2	25	2806	12	
11822	1	14	4435	11	
8667	2	16	7950	10	
10825	1	3	7144	5	

	IncidentAddress_freq
20936	28
19307	16
2609	39
15420	15
21007	21
...	...
27962	24
10263	22
11822	22
8667	46
10825	14

[23068 rows x 17 columns]

#catergorical columns

```
df_cat = X_train_filled.select_dtypes(include=['object', 'category'])
df_cat
```

	TypeOfIncident	TypeOfCollission	SeverityOfIncident	\
20936	Multi-vehicle Collision	Front Collision	Major Damage	
19307	Multi-vehicle Collision	Side Collision	Minor Damage	
2609	Vehicle Theft	NaN	Trivial Damage	
15420	Multi-vehicle Collision	Side Collision	Total Loss	
21007	Single Vehicle Collision	Side Collision	Minor Damage	
...	
27962	Multi-vehicle Collision	Side Collision	Major Damage	
10263	Multi-vehicle Collision	Rear Collision	Major Damage	
11822	Multi-vehicle Collision	Front Collision	Major Damage	
8667	Multi-vehicle Collision	Front Collision	Minor Damage	
10825	Single Vehicle Collision	Rear Collision	Minor Damage	

	AuthoritiesContacted	IncidentState	IncidentCity	PropertyDamage	\
20936	Fire	State4	City4	NaN	

19307	Other	State8	City7	NaN
2609	None	State9	City2	NaN
15420	Other	State5	City7	NaN
21007	Fire	State7	City4	NO
...
27962	Other	State5	City5	NO
10263	Police	State5	City2	YES
11822	Other	State4	City3	NO
8667	Police	State7	City5	NO
10825	Ambulance	State9	City2	YES
	Witnesses	PoliceReport	InsuredGender	InsuredEducationLevel \
20936	1	YES	MALE	Masters
19307	0	NaN	MALE	MD
2609	2	NO	MALE	Masters
15420	0	NO	MALE	JD
21007	0	NaN	FEMALE	MD
...
27962	3	NO	MALE	MD
10263	0	NO	FEMALE	College
11822	3	NO	MALE	MD
8667	2	YES	MALE	College
10825	2	NaN	MALE	JD
	InsuredOccupation	InsuredHobbies	InsurancePolicyState \	
20936	prof-specialty	movies	State1	
19307	protective-serv	sleeping	State2	
2609	transport-moving	video-games	State1	
15420	exec-managerial	paintball	State1	
21007	sales	kayaking	State3	
...	
27962	machine-op-inspct	camping	State3	
10263	craft-repair	skydiving	State3	
11822	priv-house-serv	board-games	State3	
8667	transport-moving	dancing	State2	
10825	prof-specialty	kayaking	State1	
	Policy_CombinedSingleLimit	InsuredRelationship	VehicleMake \	
20936	100/300	husband	Mercedes	
19307	100/300	own-child	Saab	

2609	500/1000	husband	Suburu
15420	250/500	other-relative	BMW
21007	100/300	not-in-family	Ford
...
27962	250/500	wife	BMW
10263	100/300	husband	Mercedes
11822	100/300	own-child	Dodge
8667	100/500	not-in-family	Audi
10825	100/300	not-in-family	BMW

	VehicleModel	Region
20936	Grand Cherokee	Tamil Nadu
19307	93	Maharashtra
2609	Forrestor	Tamil Nadu
15420	X5	Maharashtra
21007	Neon	Tamil Nadu
...
27962	M5	Madhya Pradesh
10263	E400	Tamil Nadu
11822	RAM	Tamil Nadu
8667	A3	Tamil Nadu
10825	X5	Maharashtra

[23068 rows x 19 columns]

```
distinct_values = df_cat.nunique()
distinct_values
```

TypeOfIncident	4
TypeOfCollission	3
SeverityOfIncident	4
AuthoritiesContacted	4
IncidentState	7
IncidentCity	7
PropertyDamage	2
Witnesses	5
PoliceReport	2
InsuredGender	2
InsuredEducationLevel	7
InsuredOccupation	14
InsuredHobbies	20
InsurancePolicyState	3
Policy_CombinedSingleLimit	9
InsuredRelationship	6
VehicleMake	15
VehicleModel	39
Region	3

dtype: int64

One-Hot Encoding

```
import pandas as pd

def one_hot_encode_for_fraud_analysis(df, target_col='ReportedFraud',
drop_first=False):
    """
        One-hot encodes all categorical columns (object/category) in the
        DataFrame
        except for the specified target column, typically used in fraud
        analysis.

        Parameters

        df : pd.DataFrame
            The DataFrame containing both numeric and categorical columns.
        target_col : str, optional
            The name of the target column (e.g., 'ReportedFraud'). If this
            column
            is a string type and you do not want to one-hot encode it,
            specify it here.
            Default is 'ReportedFraud'.
        drop_first : bool, optional
            Whether to drop the first category in each encoded feature to
            avoid
            the dummy variable trap. Default is False.

        Returns

        pd.DataFrame
            A new DataFrame with the categorical features encoded except
            for the target column.

    """
    # Identify all object/category columns
    cat_cols = df.select_dtypes(include=["object",
"category"]).columns.tolist()

    # Exclude the target column if in cat_cols
    if target_col in cat_cols:
        cat_cols.remove(target_col)

    # One-hot encode the remaining categorical columns
    df_encoded = pd.get_dummies(df, columns=cat_cols,
drop_first=drop_first)

    return df_encoded

X_train_ohe =
one_hot_encode_for_fraud_analysis(X_train_filled,drop_first=True)
```

```

X_test_ohe =
one_hot_encode_for_fraud_analysis(X_test_filled,drop_first=True)

print(X_train_ohe.shape)
print(X_test_ohe.shape)

(23068, 154)
(5768, 154)

```

Check Point

```

print(X_train_ohe.shape)
print(y_train.shape)

print(X_test_ohe.shape)
print(y_test.shape)

#previous results
# (92275, 40)
# (23069, 40)
# (92275,)
# (23069,)

(23068, 154)
(23068,)
(5768, 154)
(5768,)

```

Numbers are as expected. I can save the preprocessed data to a checkpoint file to avoid repeating the preprocessing steps in case of any issues. This will allow me to reload the preprocessed data and continue with the analysis without starting from scratch.

```

#checkpoint

X_train_ohe.to_csv('checkpoint/encoded_xtrain_checkpoint.csv', index =
False)
X_test_ohe.to_csv('checkpoint/encoded_xtest_checkpoint.csv', index =
False)

y_train.to_csv('checkpoint/encoded_ytrain_checkpoint.csv', index =
False)
y_test.to_csv('checkpoint/encoded_ytest_checkpoint.csv', index =
False)

```

4. Technique 1: P (Multi-Layer Perceptron)

4.1 Motivation for choosing the technique and schematic figure of the analysis process

```
from IPython.display import Image, display
import os
```

Rationale for Choosing a Multi-Layer Perceptron (MLP) for Binary Classification

I had tried a number of different models and found the MLP gave one of the best results. A Multi-Layer Perceptron (MLP) is well-suited for binary classification tasks because it can model complex, non-linear relationships between features and the target variable. By utilising multiple hidden layers and non-linear activation functions, MLPs can learn intricate decision boundaries that would be difficult to achieve with linear models alone.

MLPs also offer a high degree of flexibility and adaptability, allowing them to capture a wide range of data patterns without the need for extensive feature engineering. As a result, they can handle large datasets and high-dimensional feature spaces effectively, can be useful for fraud detection, where subtle anomalies may not follow simple patterns.

Building pipeline for the MLP

The pipeline consists of the following key components:

- Pipeline: built with scikit-learn
- Data Scaling (StandardScaler): normalises features to a mean of 0 and standard deviation of 1, helping the neural network converge faster by ensuring features are on a similar scale.
- Feature Selection (SelectKBest): uses an ANOVA F-test to remove less informative features, allowing the MLP to focus on the 20 most relevant inputs.
- Classifier (MLPClassifier): A Multi-Layer Perceptron with hidden layers of neurons and a ReLU activation function (see the next section for details).
- Early Stopping: prevents overfitting and saves time by halting training when the validation loss stops decreasing, capturing the minimal point on the learning curve.
 - Class Imbalance: as the MLP does not have built-in handling for class imbalance, the `imblearn` library (part of the scikit-learn contrib projects) is used.
 - SMOTE: creates synthetic samples of the minority class.
 - ClusterCentroids: undersamples the majority class.

MLP Scheme

The pipeline ensures the same scaling and feature selection steps are applied in every cross-validation fold during hyperparameter tuning. StratifiedKFold splits the data in a way that keeps class proportions the same, making the results fair. Finally, to handle uneven class sizes, we calculate class weights and use them in training so that the minority class is not overshadowed by the majority class.

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.feature_selection import RFE
from sklearn.decomposition import PCA, TruncatedSVD
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

y_train_mapped = y_train.map({'Y': 1, 'N': 0})
y_train_mapped.value_counts()

ReportedFraud
0    16840
1     6228
Name: count, dtype: int64

#map the target column to 1 and 0

y_test_mapped = y_test.map({'Y': 1, 'N': 0})
y_test_mapped.value_counts()

ReportedFraud
0     4211
1     1557
Name: count, dtype: int64

print(X_train_ohe.shape)
print(y_train_mapped.shape)

print(X_test_ohe.shape)
print(y_test_mapped.shape)

#previous results
# (92275, 40)
# (23069, 40)
# (92275,)
# (23069,)

(23068, 154)
(23068,)
(5768, 154)
(5768,)

X_train = X_train_ohe
y_train = y_train_mapped
```



```

X_test = X_test_ohe
y_test = y_test_mapped

import numpy as np
from sklearn.model_selection import train_test_split, StratifiedKFold,
GridSearchCV, cross_val_predict
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.utils.class_weight import compute_class_weight

```

4.2 Other items necessary for the technique

```

#pipeline
# Part of scikit learn
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline
from imblearn.under_sampling import ClusterCentroids # for
undersampling to improve class balance

pipeline = ImbPipeline([
    ('smote', SMOTE(random_state=42)), # Add SMOTE step
    #('undersample', ClusterCentroids(random_state=42)), #
undersampling step
    ('scaler', StandardScaler()), # Scale features to have mean=0 and
variance=1
    ('feat_select', SelectKBest(score_func=f_classif, k=20)), # Select
the 20 best features based on ANOVA F-value
    ('clf', MLPClassifier( # Multi-layer Perceptron classifier
        random_state=42, # for reproducibility
        early_stopping=True, # stop training when performance worsens
on the validation set
        validation_fraction=0.2, # use 20% of the training data as
validation
        n_iter_no_change=5 # stop training if the validation score
doesn't improve for 5 iterations
    ))
])

```

4.3 Setting hyperparameters (rationale)

- param_grid: Set up to explore a range of hyperparameter values to optimise the performance of the Multi-Layer Perceptron (MLP) classifier within the pipeline. The chosen hyperparameters and their respective values are justified as follows:
- feat_select_k: The number of top features to select using SelectKBest. The step will try 10, or 20 or 30 features, and see which works best. Testing these values allows

for identifying the subset of features that contribute most to the predictive power, reducing the risk of overfitting and improving computational efficiency.

- `hidden_layer_sizes`: This is the MLPClassifier (neural network) step. Each tuple represents the number of neural networks hidden layer; initially set to: (100,) - one hidden layer with 100 neurons. (100,50) - 2 hidden layers, one with 100 neurons then another with 50 neurons. (100,100) means 2 hidden layers, each with 100 neurons. These can be adjusted to create different shapes of the network to see which performs best. I will use a 3-layer architecture which is deep enough to capture complex patterns in the data without being too deep to cause overfitting.
- `Regulisation: alpha (L2 regularisation)`: The regularisation parameter to prevent overfitting and penalising large weights. I will set this to [0.0001, 0.005, 0.01] to a weaker regularisation so it can learn more complex patterns but may overfit. Exploring these alpha values allows for finding the right balance between fitting the training data and maintaining generalization to unseen data.
- `learning_rate_init`: The initial learning rate for is [0.001, 0.005, 0.01]. A higher rate can make the model faster but may overshoot the best solution; while a lower rate can help the model learn more carefully but may take longer to train.

By systematically tuning these hyperparameters, we aim to enhance the model's ability to generalise, achieve better performance metrics, and ensure robustness in fraud detection tasks.

```
param_grid = {
    'feat_select_k': [10, 20, 30], # Number of features to select
    'clf_hidden_layer_sizes': [(100,), (100, 50), (100, 100)], #
    Number of neurons in each hidden layer
    'clf_alpha': [0.001, 0.01, 0.5], # L2 penalty (regularisation
    term) parameter
    'clf_learning_rate_init': [0.001, 0.005, 0.01], # Initial
    learning rate
}
```

4.4 Dealing with class imbalance, if necessary, if not done above

```
# percentage of each class in y_train
fraud_train_counts = y_train.value_counts()
fraud_train_percent = y_train.value_counts(normalize=True) * 100
fraud_train_percent
```

```
ReportedFraud
0    73.001561
1    26.998439
Name: proportion, dtype: float64
```

Balanced class weights are needed to prevent the model from being biased towards the majority class. This is especially important in fraud detection tasks where the number of fraudulent claims is much lower than genuine claims. To address this issue, I have used `imblearn` in the

pipeline as MLP doesn't have built-in handling for class imbalance and neither does a manual entry of class weights work.

```
# Code moved to pipeline
```

4.5 Optimising the Hyperparameters Appropriately

- **Cross-Validation Strategy:** stratified k-fold cross-validation with 3 splits to ensure that each fold maintains the class distribution of the original data. This will help to evaluate the model's performance across different subsets of the data and provide a more reliable estimate.
- **GridSearchCV with MLPClassifier:** Grid search with cross-validation will find the optimal hyperparameters for the MLP classifier by trying different combinations from the `param_grid` on the given pipeline. The accuracy metric is used to measure performance and cross validation is defined by the `cv_strategy`. Using this over the parameter grid, the model will be trained and evaluated on different hyperparameter combinations to identify the best configuration that maximises performance metrics. With `n_jobs=-1`, all CPU cores will be used to speed up the search.
- **Grid_search.fit(X_train, y_train)** splits the training data, trains the pipeline with every hyperparameter combination, measures accuracy, and selects the best-performing set of parameters.

[illegible]

```

        ('feat_select',
SelectKBest(k=20)),
        ('clf',
MLPClassifier(early_stopping=True,
n_iter_no_change=5,
random_state=42,
validation_fraction=0.2))),
        n_jobs=-1,
        param_grid={'clf__alpha': [0.001, 0.01, 0.5],
                    'clf__hidden_layer_sizes': [(100,), (100,
50),
                    (100, 100)],
                    'clf__learning_rate_init': [0.001, 0.005,
0.01],
                    'feat_select__k': [10, 20, 30]},
                    scoring='accuracy', verbose=1)

```

4.5 Performance Metrics for Training Set

- Best parameters from grid search, shows the best cross-validated score, and the best estimator from the grid search. The best estimator is the model with the optimal hyperparameters that achieved the highest cross-validated score during grid search.
- Cross val prediction score obtains out of the box predictions from the cross-validated model. This allows us to evaluate the model's performance on the training set using the best hyperparameters identified during grid search.

```

print("Best Params:", grid_search.best_params_)
print("Best CV Score on Training Splits:", grid_search.best_score_)

y_val_pred = cross_val_predict(
    grid_search.best_estimator_, # The best pipeline from GridSearchCV
    X_train, # Training data
    y_train, # Target labels
    cv=cv_strategy, # Same CV strategy as GridSearchCV
    n_jobs=-1
)

best_model_mlp = grid_search.best_estimator_

accuracy = accuracy_score(y_train, y_val_pred) # Compute overall
accuracy
print(f"K-Fold Averaged Accuracy: {accuracy:.4f}")

class_report = classification_report(y_train, y_val_pred) # Generate

```

```

classification report
print("\nClassification Report:\n", class_report)

Best Params: {'clf__alpha': 0.01, 'clf__hidden_layer_sizes': (100, 100), 'clf__learning_rate_init': 0.005, 'feat_select__k': 30}
Best CV Score on Training Splits: 0.826512618891663
K-Fold Averaged Accuracy: 0.8265

```

Classification Report:				
	precision	recall	f1-score	support
0	0.87	0.90	0.88	16840
1	0.70	0.62	0.66	6228
accuracy			0.83	23068
macro avg	0.78	0.76	0.77	23068
weighted avg	0.82	0.83	0.82	23068

Classification Report from Training Set MLP

Best Hyperparameters: -clf__alpha': 0.01 - moderate level of L2 regularisation on MLP

- clf__hidden_layer_sizes': (100, 100) - two hidden layers with 100 neurons each
- 'clf__learning_rate_init': 0.005, - moderate learning rate
- 'feat_select__k': 30 - 30 top features selected by SelectKBest

Best Cross-Validated Score: 0.8265

- On average, across multiple train/validation splits, this configuration correctly predicted the label about 82.65% of the time.

K-Fold Cross-Validation Score: 0.8265

- The model achieved an average accuracy of 82.65% across the training data using the best hyperparameters identified during grid search.

Classification Report:

- class 0 - precision = 0.87, recall = 0.90, f1-score = 0.88; the model is good at identifying Class 0 with predicts with correct 87% of the time and recalls 89% of the time.
- class 1 - precision = 0.70, recall = 0.62, f1-score = 0.66; the model is not as good at identifying Class 1 with predicts with correct 70% of the time and recalls 62% of the time.
- Overall Accuracy across both classes is 83%
- Weighted average (0.82 precision, 0.83 recall) accounts for the fact that Class 0 has more samples and thus carries more weight in the overall score.

- Overall the model achieves around 83% accuracy on the training dataset, performs very well on the majority class (Class 0), and is reasonably effective but less accurate for the minority class (Class 1).
- ('undersample', ClusterCentroids(random_state=42)), undersampling step was removed as it did not improve performance and also significantly increased the processing it.

5. Technique 2 - HistGradientBoostClassifier

5.1 Motivation for choosing the technique and schematic figure of the analysis process

For the binary classification problem, I chose a tree-based gradient boosting approach and which uses an ensemble method. The HistGradientBoostingClassifier is preferred over the GradientBoostingClassifier, due to its efficiency of speed as it uses bins to group continuous features. It is also more memory efficient, is robust to outliers and can handle imbalanced datasets effectively. Scaling is not required for decision tree-based models, making it computationally efficient for large datasets. This makes it well-suited for binary classification tasks as it can handle complex relationships between features and the target variable, making it effective for fraud detection.

The schematic analysis flow is as follows:

Histogram Gradient Boosting Classifier

5.2 Other items necessary for the technique

Pipeline

```
from sklearn.pipeline import Pipeline
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.model_selection import StratifiedKFold, GridSearchCV,
cross_val_predict
from sklearn.metrics import accuracy_score, classification_report

# 1) pipeline feature selection HistGradientBoosting
pipeline = Pipeline([
    ('feat_select', SelectKBest(score_func=f_classif, k=20)),
    ('clf', HistGradientBoostingClassifier( # Gradient Boosting
        random_state=42, # for reproducibility
        early_stopping=True, # built-in early stopping
        validation_fraction=0.2, # use 20% of the training data as
validation set
        n_iter_no_change=5 # stop training if the validation
score doesn't improve for 5 iterations
```

```
))  
])
```

5.3 Setting hyperparameters (rationale)

The hyperparameters for the HistGradientBoostingClassifier are chosen based on the following rationale:

- 'feat_select_k': [10, 20, 30] - I am looking for the top features to keep after the F-test in SelectKBest. I will test different values to identify the subset of features that contribute most to the predictive power. The aim is to find the sweet spot too many features can lead to overfitting, while too few features can lead to underfitting.
- clf_max_iter': [100, 200] - The maximum number of iterations for the gradient boosting classifier. I will test different values to find the optimal number of iterations for convergence. The aim is calancing compute time with thorough training. Higher accuracy can be achieved with more iterations, but it comes at the cost of increased computational time and potential overfitting.
- clf_learning_rate': [0.05, 0.1, 0.2] - The learning rate shrinks the contribution of each tree. I will test different values to control the impact of each tree on the final prediction. Lower valus can give better training, while higher values can be faster but risk overfitting.
- clf_max_depth': [None, 5] - The maximum depth of the decision trees. I will test different values to control the complexity of the trees and prevent overfitting. A limit of 5 can help prevent overfitting.
- clf_max_leaf_nodes': [31, 63] - The maximum number of leaf nodes in each tree. I will test different values to control the complexity of the trees and prevent overfitting. A limit of 31 can help prevent overfitting. A larger leaf count can lead to more complex trees and potential overfitting and become computationally inefficient.

```
# 2) hyperparameter search space  
param_grid = {  
    'feat_select_k': [10, 20, 30], # Number of features to select  
    'clf_max_iter': [100, 200], # Maximum number of iterations  
    'clf_learning_rate': [0.05, 0.1, 0.2], # Learning rate  
    'clf_max_depth': [None, 5], # Maximum depth of the trees  
    'clf_max_leaf_nodes': [31, 63], # Maximum number of leaf nodes  
    per tree  
}
```

5.4 Optimising the Hyperparameters Appropriately

Cross Validation Strategy: The training data is split into 3 folds (subsets). Stratified makes each fold preserves the proportion of each class. Thhe data is shuffled before splitting to make the folds more random. Setting a random_state (42) makes sure results are reproducible—if someone else runs the code, they'll get the same splits. The best hyperparameters are chosen based on the average score across the cross-validation folds.

GridSearchCV is used over the specified paramater grid to find the optimal hyperparameters for the HistGradientBoostingClassifier to tune the pipeline. Performance is measured by accuracy. . The grid search will train and evaluate the model on different hyperparameter combinations to identify the best configuration that maximises performance metrics. The final model is evaluated on the train set to assess its performance.

```
cv_strategy = StratifiedKFold(n_splits=3, shuffle=True,
random_state=42)
grid_search = GridSearchCV( # Grid search for hyperparameter tuning
    estimator=pipeline, # The pipeline to tune
    param_grid=param_grid, # The hyperparameter search space
    scoring='accuracy', # The metric to optimise
    cv=cv_strategy, # The cross-validation strategy
    n_jobs=-1, # Use all available CPU cores
    verbose=1 # Output progress
)
grid_search.fit(X_train, y_train)
```

Fitting 3 folds for each of 72 candidates, totalling 216 fits

```
GridSearchCV(cv=StratifiedKFold(n_splits=3, random_state=42,
                                shuffle=True),
              estimator=Pipeline(steps=[('feat_select',
                                         SelectKBest(k=20)),
                                         ('clf',
                                          HistGradientBoostingClassifier(early_stopping=True,
                                                                           n_iter_no_change=5,
                                                                           random_state=42,
                                                                           validation_fraction=0.2))]),
              n_jobs=-1,
              param_grid={'clf__learning_rate': [0.05, 0.1, 0.2],
                          'clf__max_depth': [None, 5],
                          'clf__max_iter': [100, 200],
                          'clf__max_leaf_nodes': [31, 63],
                          'feat_select__k': [10, 20, 30]},
              scoring='accuracy', verbose=1)
```


5.5 Performance metrics for training set

```
print("Best Params:", grid_search.best_params_)
print("Best CV Score on Training Splits:", grid_search.best_score_)

y_val_pred = cross_val_predict(
    grid_search.best_estimator_, # The best pipeline from GridSearchCV
    X_train, # Training data
    y_train, # Target labels
    cv=cv_strategy, # Same CV strategy as GridSearchCV
    n_jobs=-1
)
```

```
best_model_hgb = grid_search.best_estimator_
```

```
accuracy = accuracy_score(y_train, y_val_pred) # Compute overall
accuracy
```

```
print(f"K-Fold Averaged Accuracy: {accuracy:.4f}")
```

```
class_report = classification_report(y_train, y_val_pred) # Generate
classification report
```

```
print("\nClassification Report:\n", class_report)
```

```
Best Params: {'clf__learning_rate': 0.1, 'clf__max_depth': None,
'clf__max_iter': 200, 'clf__max_leaf_nodes': 63, 'feat_select__k': 30}
Best CV Score on Training Splits: 0.8755854080973934
K-Fold Averaged Accuracy: 0.8756
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.89         0.94         0.92       16840
     1       0.82         0.70         0.75        6228

 accuracy                   0.88       23068
 macro avg              0.85         0.82         0.83       23068
 weighted avg           0.87         0.88         0.87       23068
```

Classification Report for Training Set HGB

1. Best Hyperparameters

- `clf__learning_rate = 0.1`: The model's learning rate is 0.1, balancing how quickly it learns.
- `clf__max_depth = None`: The model's trees can grow without a predefined maximum depth (they can become quite deep if needed).
- `clf__max_iter = 200`: Up to 200 boosting stages (iterations) are allowed.

- `clf__max_leaf_nodes = 63`: Each tree can have at most 63 leaf nodes.
 - `feat_select__k = 30`: The model is using the top 30 features, as selected by `SelectKBest`.
2. **Best Cross-Validation Score**
- 0.875585... (about 87.56%).
 - This means that, on average, across the cross-validation folds, the model correctly classified about 87.56% of instances with the best hyperparameter settings.
3. **K-Fold Averaged Accuracy**
- 0.8756 (nearly 87.56%).
 - When we look at the predictions over the entire training set using cross-validation, the model is again correct about 87.56% of the time, which is consistent with the best CV score.
4. **Classification Report**
- Overall Accuracy: 0.88 or 88% (slightly rounded).
 - Class 0:
 - Precision: 0.89 (when predicting class 0, it's right 89% of the time).
 - Recall: 0.94 (out of all actual class 0 samples, it correctly identifies 94%).
 - F1-score: 0.92 (balanced measure of precision and recall).
 - Class 1:
 - Precision: 0.82 (when predicting class 1, it's right 82% of the time).
 - Recall: 0.70 (it catches 70% of the true class 1 instances).
 - F1-score: 0.75.
 - The model clearly performs better on class 0 (the majority class) than on class 1, shown by higher recall and precision for class 0.
 - Macro Average (0.85 precision, 0.82 recall, 0.83 F1) treats both classes equally.
 - Weighted Average (0.87 precision, 0.88 recall, 0.87 F1) accounts for the fact that there are more class 0 samples, so the overall score is skewed toward class 0's strong performance.

Summary:

- The model is quite accurate overall (~88%).

- It classifies class 0 more confidently and more often correctly than class 1. This is typical when the dataset is imbalanced, but still, an F1-score of 0.75 for class 1 indicates a reasonably good performance for the minority class.

1. MLPClassifier Performance

Test Accuracy: 0.860957

Balanced Accuracy: 0.800743

Classification Report

Class	Precision	Recall	F1-score	Support
N	0.88	0.93	0.91	4211
Y	0.78	0.67	0.72	1557
Accuracy			0.86	5768
Macro Avg	0.83	0.80	0.81	5768
Weighted Avg	0.86	0.86	0.86	5768

Confusion Matrix

	Predicted N	Predicted Y
Actual N	3923	288
Actual Y	514	1043

2. HistGradientBoostingClassifier Performance

Test Accuracy: 0.868412

Balanced Accuracy: 0.814551

Classification Report

Class	Precision	Recall	F1-score	Support
N	0.89	0.93	0.91	4211
Y	0.79	0.70	0.74	1557
Accuracy			0.87	5768
Macro Avg	0.84	0.81	0.83	5768
Weighted Avg	0.87	0.87	0.87	5768

Confusion Matrix

	Predicted N	Predicted Y
Actual N	3923	288
Actual Y	471	1086

Summary Comparison

Metric	MLPClassifier	HistGradientBoostingClassifier
Test Accuracy	86.10%	86.84%
Balanced Accuracy	80.07%	81.46%
Precision (N)	0.88	0.89
Precision (Y)	0.78	0.79
Recall (N)	0.93	0.93
Recall (Y)	0.67	0.70
F1-score (N)	0.91	0.91
F1-score (Y)	0.72	0.74
Support (N)	4211	4211
Support (Y)	1557	1557
Overall Accuracy	0.86	0.87

Observations

- **Test Accuracy:** Both models perform similarly, with the HistGradientBoostingClassifier having a slight edge.
- **Balanced Accuracy:** HistGradientBoostingClassifier shows better balanced accuracy, indicating improved performance on minority classes.
- **Precision & Recall:** Both models have high precision and recall for class 'N' ("No Fraud"). The Gradient Boosting model slightly outperforms the MLPClassifier in precision and recall for class 'Y' ("Fraud").
- **F1-score:** The Gradient Boosting model has a marginally higher F1-score for class 'Y', suggesting better balance between precision and recall.
- **Confusion Matrix:** The Gradient Boosting model correctly identifies more fraud cases (higher true positives) while maintaining a similar number of false positives compared to the MLPClassifier.

These tables provide a comprehensive overview of the performance metrics for both classifiers, facilitating an informed comparison to determine which model better suits your objectives, especially concerning fraud detection.

6 Comparing the Two Techniques

- Use of nested cross-validation for both techniques to deal with overfitting model selection and model comparison
- Use appropriate metrics for the testing set
- Use an appropriate model selection visualisation curve (ROC, PR, etc.) that is suitable for the problem at hand
- Checking for overfitting

6.1 Use of Nested Cross-Validation

I will use set up a nested cross-validation to evaluate the performance of both models on the test set. This will help to ensure that the models are robust and generalise well to unseen data. I will use the same metrics as before to evaluate the models and compare their performance. This will include:

1. Two pipelines for the MLPClassifier and HistGradientBoostingClassifier.
2. Hyperparameter tuning using GridSearchCV.
3. Define Outer CV (5 folds) and Inner CV (3 folds) for nested cross-validation.
4. Split data into train_outer and test_outer
5. Inner CV: Grid search over hyperparameters for each model.
6. Evaluate the best models from each technique on the test_outer set.
7. Calculate and compare performance metrics for both models on the test set.

```
import numpy as np
from sklearn.model_selection import StratifiedKFold,
RandomizedSearchCV
from sklearn.pipeline import Pipeline
from sklearn.neural_network import MLPClassifier
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.metrics import accuracy_score
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline
from sklearn.ensemble import HistGradientBoostingClassifier

# Example pipeline for HGB: feature selection +
HistGradientBoostingClassifier
pipeline_hgb = Pipeline([
    ('feat_select', SelectKBest(score_func=f_classif)),
    ('clf', HistGradientBoostingClassifier(random_state=42,
early_stopping=True, validation_fraction=0.1))
])

# Example hyperparameter distributions for HGB
param_distributions_hgb = {
    'feat_select__k': [10, 20, 30],
    'clf__max_iter': [100, 200],
    'clf__learning_rate': [0.05, 0.1, 0.2],
    'clf__max_depth': [None, 5],
    'clf__max_leaf_nodes': [31, 63]
}

# Outer cross-validation strategy (for model evaluation)
outer_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Inner cross-validation strategy (for hyperparameter tuning)
inner_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```

# Initialize lists to store outer scores and best models
outer_scores = [] # To store scores of the outer loop
best_models = [] # To store best models for each fold

# Nested cross-validation: iterate through outer folds
for train_idx, val_idx in outer_cv.split(X_train, y_train):
    # Split data into train and test for the outer loop
    X_train_curr, X_val_curr = X_train.iloc[train_idx],
X_train.iloc[val_idx]
    y_train_curr, y_val_curr = y_train.iloc[train_idx],
y_train.iloc[val_idx]

    # Randomized search in the inner loop for hyperparameter tuning
    grid_search = RandomizedSearchCV(
        estimator=pipeline_hgb,
        param_distributions=param_distributions_hgb,
        scoring='accuracy',
        cv=inner_cv,
        n_jobs=-1,
        verbose=1
    )

    # Fit the grid search on the inner training set
    grid_search.fit(X_train_curr, y_train_curr)

    # Evaluate the best model on the outer validation set
    best_model = grid_search.best_estimator_
    y_pred = best_model.predict(X_val_curr)
    outer_score = accuracy_score(y_val_curr, y_pred)

    # Append outer score and model to the lists
    outer_scores.append(outer_score)
    best_models.append(best_model)

    print(f"Outer fold score: {outer_score:.4f}")
    print(f"Best hyperparameters for this fold:
{grid_search.best_params_}")

# Report final nested cross-validation score
print(f"Nested CV Accuracy: {np.mean(outer_scores):.4f} ±
{np.std(outer_scores):.4f}")

# Identify the best model based on outer fold performance
best_outer_fold_idx = np.argmax(outer_scores)
best_model_from_outer_hgb = best_models[best_outer_fold_idx]

print(f"Best model (based on outer fold score):
{best_model_from_outer_hgb}")
print(f"Best outer fold score:
{outer_scores[best_outer_fold_idx]:.4f}")

```

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8825
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__max_leaf_nodes': 63, 'clf__max_iter': 200, 'clf__max_depth':
None, 'clf__learning_rate': 0.05}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8680
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__max_leaf_nodes': 63, 'clf__max_iter': 200, 'clf__max_depth': 5,
'clf__learning_rate': 0.2}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8723
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__max_leaf_nodes': 31, 'clf__max_iter': 200, 'clf__max_depth':
None, 'clf__learning_rate': 0.1}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8762
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__max_leaf_nodes': 63, 'clf__max_iter': 200, 'clf__max_depth':
None, 'clf__learning_rate': 0.2}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8760
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__max_leaf_nodes': 63, 'clf__max_iter': 200, 'clf__max_depth': 5,
'clf__learning_rate': 0.1}
Nested CV Accuracy: 0.8750 ± 0.0048
Best model (based on outer fold score):
Pipeline(steps=[('feat_select', SelectKBest(k=30)),
                ('clf',
                 HistGradientBoostingClassifier(early_stopping=True,
                                                  learning_rate=0.05,
                                                  max_iter=200,
                                                  max_leaf_nodes=63,
                                                  random_state=42))])
Best outer fold score: 0.8825

# Example: Pipeline for MLP
pipeline_mlp = ImbPipeline([
    ('smote', SMOTE(random_state=42)), # Add SMOTE step
    #('undersample', ClusterCentroids(random_state=42)), #
undersampling step
    ('scaler', StandardScaler()), # Scale features to have mean=0 and
variance=1
    ('feat_select', SelectKBest(score_func=f_classif)),
    ('clf', MLPClassifier(random_state=42, early_stopping=True,
validation_fraction=0.1))
])

# Hyperparameter distributions for MLP (renamed from 'param_grid_mlp')

```

```

param_distributions_mlp = {
    'feat_select_k': [10, 20, 30], # Number of features to select
    'clf_hidden_layer_sizes': [(100, ), (100, 50), (100, 100)],
    'clf_alpha': [0.001, 0.01, 0.5],
    'clf_learning_rate_init': [0.001, 0.005, 0.01]
}

# Outer cross-validation strategy (for model evaluation)
outer_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Inner cross-validation strategy (for hyperparameter tuning)
inner_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Initialize lists to store outer scores and best models
outer_scores = [] # To store scores of the outer loop
best_models = [] # To store best models for each fold

# Nested cross-validation: iterate through outer folds
for train_idx, val_idx in outer_cv.split(X_train, y_train):
    # Split data into train and test for the outer loop
    X_train_curr, X_val_curr = X_train.iloc[train_idx],
X_train.iloc[val_idx]
    y_train_curr, y_val_curr = y_train.iloc[train_idx],
y_train.iloc[val_idx]

    # Randomized search in the inner loop for hyperparameter tuning
    grid_search = RandomizedSearchCV(
        estimator=pipeline_mlp,
        param_distributions=param_distributions_mlp,
        scoring='accuracy',
        cv=inner_cv,
        n_jobs=-1,
        verbose=1
    )

    # Fit the grid search on the inner training set
    grid_search.fit(X_train_curr, y_train_curr)

    # Evaluate the best model on the outer validation set
    best_model = grid_search.best_estimator_
    y_pred = best_model.predict(X_val_curr)
    outer_score = accuracy_score(y_val_curr, y_pred)

    # Append outer score and model to the lists
    outer_scores.append(outer_score)
    best_models.append(best_model)

    print(f"Outer fold score: {outer_score:.4f}")
    print(f"Best hyperparameters for this fold:
{grid_search.best_params_}")

```



```

# Report final nested cross-validation score
print(f"Nested CV Accuracy: {np.mean(outer_scores):.4f} ±
{np.std(outer_scores):.4f}")

# Identify the best model based on outer fold performance
best_outer_fold_idx = np.argmax(outer_scores)
best_model_from_outer_mlp = best_models[best_outer_fold_idx]

print(f"Best model (based on outer fold score):
{best_model_from_outer_mlp}")
print(f"Best outer fold score:
{outer_scores[best_outer_fold_idx]:.4f}")

```

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8251
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__learning_rate_init': 0.01, 'clf__hidden_layer_sizes': (100, 50),
'clf__alpha': 0.01}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8244
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__learning_rate_init': 0.01, 'clf__hidden_layer_sizes': (100, 50),
'clf__alpha': 0.01}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8147
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__learning_rate_init': 0.01, 'clf__hidden_layer_sizes': (100, 50),
'clf__alpha': 0.001}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8229
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__learning_rate_init': 0.001, 'clf__hidden_layer_sizes': (100,),
'clf__alpha': 0.01}
Fitting 5 folds for each of 10 candidates, totalling 50 fits
Outer fold score: 0.8240
Best hyperparameters for this fold: {'feat_select_k': 30,
'clf__learning_rate_init': 0.005, 'clf__hidden_layer_sizes': (100,
50), 'clf__alpha': 0.001}
Nested CV Accuracy: 0.8222 ± 0.0038
Best model (based on outer fold score): Pipeline(steps=[('smote',
SMOTE(random_state=42)), ('scaler', StandardScaler()),
('feat_select', SelectKBest(k=30)),
('clf',
MLPClassifier(alpha=0.01, early_stopping=True,
hidden_layer_sizes=(100, 50),
learning_rate_init=0.01,
random_state=42))])
Best outer fold score: 0.8251

```

6.2 Use appropriate metrics for the testing set

```
best_model_mlp = best_model_from_outer_mlp
best_model_hgb = best_model_from_outer_hgb

## Evaluation on MLP on TEST SET

y_test_pred_mlp = best_model_mlp.predict(X_test)

# Accuracy
acc = accuracy_score(y_test, y_test_pred_mlp)
print(f"Test Accuracy (MLP): {acc:.4f}")

# Classification report (precision, recall, F1, support)
print("Classification Report (MLP):")
print(classification_report(y_test, y_test_pred_mlp))

# Confusion matrix
print("Confusion Matrix (MLP):")
print(confusion_matrix(y_test, y_test_pred_mlp))
```

Test Accuracy (MLP): 0.8225

Classification Report (MLP):

	precision	recall	f1-score	support
0	0.87	0.89	0.88	4211
1	0.69	0.63	0.66	1557
accuracy			0.82	5768
macro avg	0.78	0.76	0.77	5768
weighted avg	0.82	0.82	0.82	5768

Confusion Matrix (MLP):

```
[[3761  450]
 [ 574  983]]
```

```
# Evaluation on HGB on TEST SET
```

```
y_test_pred_hgb = best_model_hgb.predict(X_test)
```

```
# Accuracy
```

```
acc = accuracy_score(y_test, y_test_pred_hgb)
print(f"Test Accuracy (HGB): {acc:.4f}")
```

```
# Classification report (precision, recall, F1, support)
```

```
print("Classification Report (HGB):")
print(classification_report(y_test, y_test_pred_hgb))
```

```
# Confusion matrix
```

```
print("Confusion Matrix (HGB):")
```

```
print(confusion_matrix(y_test, y_test_pred_hgb))
```

Test Accuracy (HGB): 0.8774

Classification Report (HGB):

	precision	recall	f1-score	support
0	0.89	0.94	0.92	4211
1	0.82	0.70	0.75	1557
accuracy			0.88	5768
macro avg	0.86	0.82	0.84	5768
weighted avg	0.87	0.88	0.87	5768

Confusion Matrix (HGB):

```
[[3978 233]
 [ 474 1083]]
```

For analysis, please see section 7.1 Technical Perspective

6.3 Appropriate Model Selection Visualisation Curves

```
from sklearn.metrics import roc_curve, auc, precision_recall_curve,
average_precision_score

y_proba_mlp = best_model_mlp.predict_proba(X_test)[: , 1]
y_proba_hgb = best_model_hgb.predict_proba(X_test)[: , 1]

# Compute ROC curve and ROC area for MLP
fpr_mlp, tpr_mlp, _ = roc_curve(y_test_mapped, y_proba_mlp)
roc_auc_mlp = auc(fpr_mlp, tpr_mlp)

# Compute ROC curve and ROC area for HistGradientBoostingClassifier
fpr_gb, tpr_gb, _ = roc_curve(y_test_mapped, y_proba_hgb)
roc_auc_gb = auc(fpr_gb, tpr_gb)

# Compute Precision-Recall curve and average precision for MLP
precision_mlp, recall_mlp, _ = precision_recall_curve(y_test_mapped,
y_proba_mlp)
avg_precision_mlp = average_precision_score(y_test_mapped,
y_proba_mlp)

# Compute Precision-Recall curve and average precision for
HistGradientBoostingClassifier
precision_gb, recall_gb, _ = precision_recall_curve(y_test_mapped,
y_proba_hgb)
```

```

avg_precision_gb = average_precision_score(y_test_mapped, y_proba_hgb)

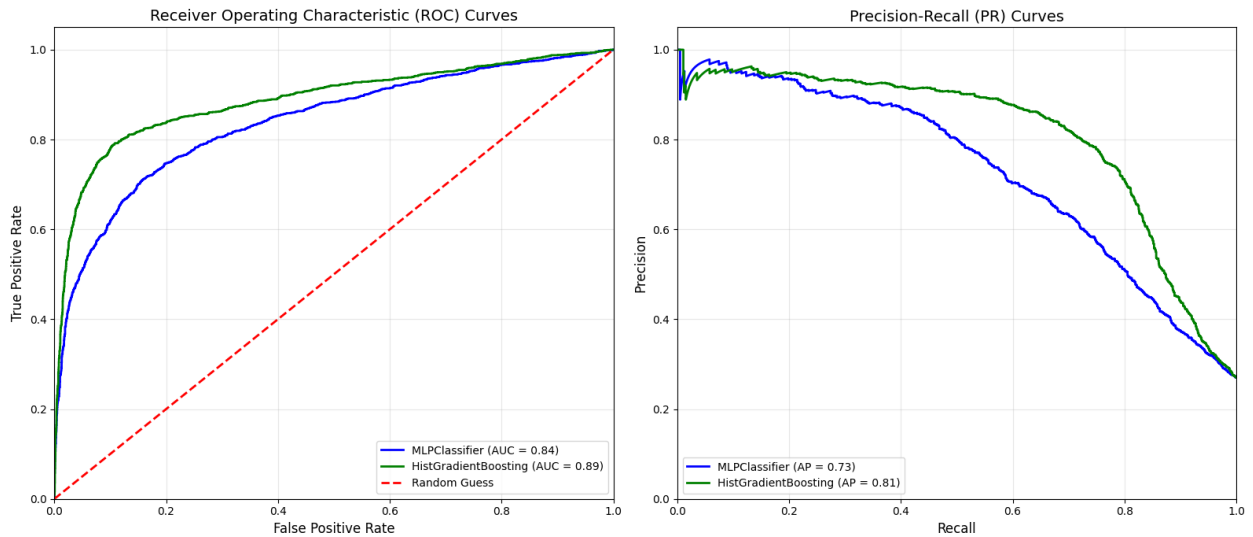
# Create subplots for ROC and Precision-Recall curves
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 7))

# Plot ROC curves
ax1.plot(fpr_mlp, tpr_mlp, color='blue', lw=2, label=f'MLPClassifier (AUC = {roc_auc_mlp:.2f})')
ax1.plot(fpr_gb, tpr_gb, color='green', lw=2, label=f'HistGradientBoosting (AUC = {roc_auc_gb:.2f})')
ax1.plot([0, 1], [0, 1], color='red', lw=2, linestyle='--', label='Random Guess')
ax1.set_xlim([0.0, 1.0])
ax1.set_ylim([0.0, 1.05])
ax1.set_xlabel('False Positive Rate', fontsize=12)
ax1.set_ylabel('True Positive Rate', fontsize=12)
ax1.set_title('Receiver Operating Characteristic (ROC) Curves', fontsize=14)
ax1.legend(loc="lower right", fontsize=10)
ax1.grid(alpha=0.3)

# Plot Precision-Recall curves
ax2.plot(recall_mlp, precision_mlp, color='blue', lw=2, label=f'MLPClassifier (AP = {avg_precision_mlp:.2f})')
ax2.plot(recall_gb, precision_gb, color='green', lw=2, label=f'HistGradientBoosting (AP = {avg_precision_gb:.2f})')
ax2.set_xlim([0.0, 1.0])
ax2.set_ylim([0.0, 1.05])
ax2.set_xlabel('Recall', fontsize=12)
ax2.set_ylabel('Precision', fontsize=12)
ax2.set_title('Precision-Recall (PR) Curves', fontsize=14)
ax2.legend(loc="lower left", fontsize=10)
ax2.grid(alpha=0.3)

plt.tight_layout()
plt.show()

```



For analysis, please see section 7.1 Technical Perspective

6.4 Checking for Overfitting

```
from sklearn.model_selection import learning_curve

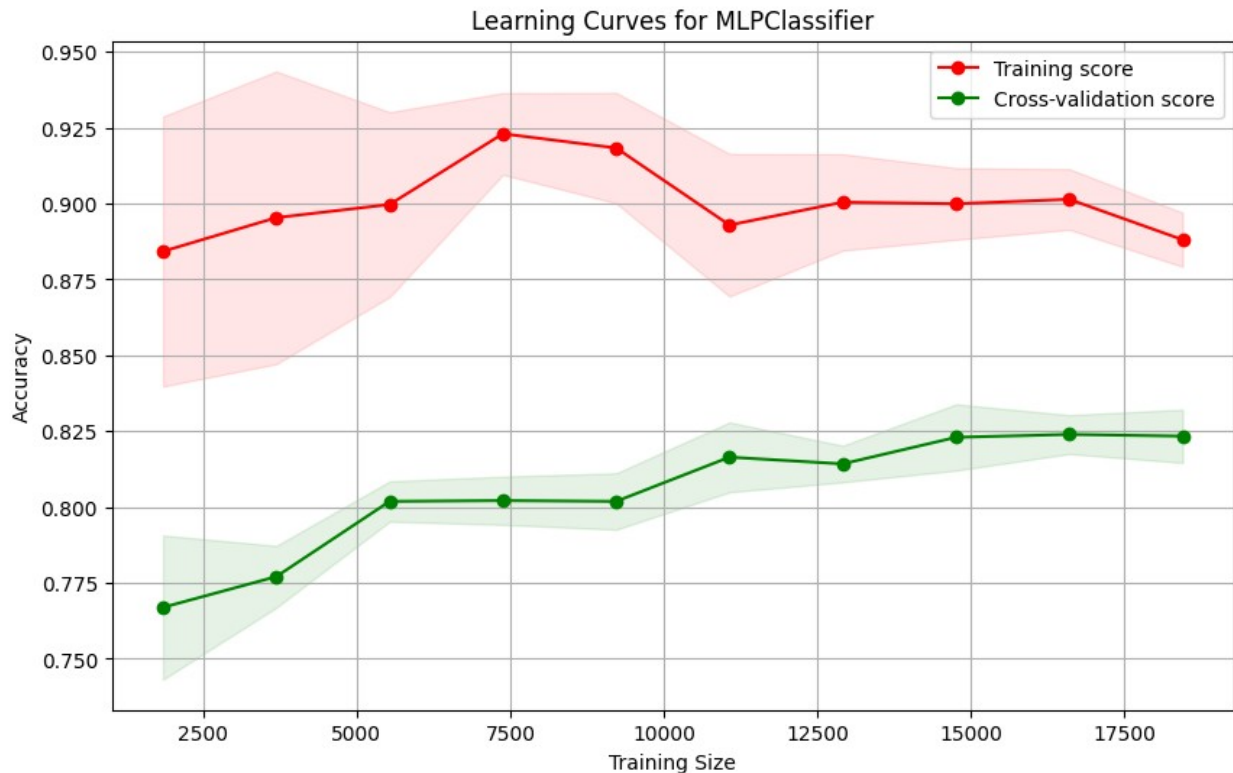
# Define cross-validation strategy
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Generate learning curve
train_sizes, train_scores, val_scores = learning_curve(
    best_model_mlp, X_train, y_train, cv=cv, scoring='accuracy',
    train_sizes=np.linspace(0.1, 1.0, 10), n_jobs=-1)

# Calculate mean and std
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
val_scores_mean = np.mean(val_scores, axis=1)
val_scores_std = np.std(val_scores, axis=1)

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_scores_mean, 'o-', color='r',
         label='Training score')
plt.plot(train_sizes, val_scores_mean, 'o-', color='g', label='Cross-
validation score')
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color='r')
plt.fill_between(train_sizes, val_scores_mean - val_scores_std,
                 val_scores_mean + val_scores_std, alpha=0.1,
                 color='g')
plt.title('Learning Curves for MLPClassifier')
```

```
plt.xlabel('Training Size')
plt.ylabel('Accuracy')
plt.legend(loc='best')
plt.grid()
plt.show()
```



```
# Define cross-validation strategy
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Generate learning curve
train_sizes, train_scores, val_scores = learning_curve(
    best_model_hgb, X_train, y_train, cv=cv, scoring='accuracy',
    train_sizes=np.linspace(0.1, 1.0, 10), n_jobs=-1)

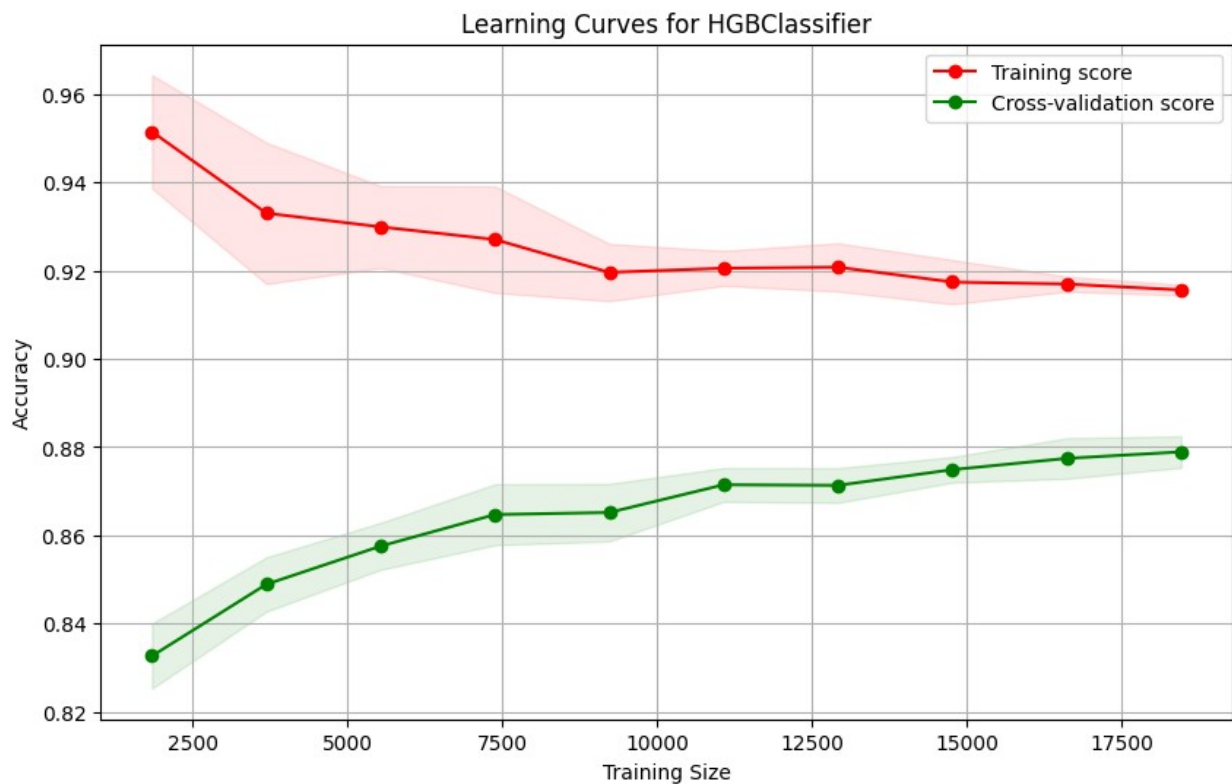
# Calculate mean and std
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
val_scores_mean = np.mean(val_scores, axis=1)
val_scores_std = np.std(val_scores, axis=1)

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_scores_mean, 'o-', color='r',
label='Training score')
```

```

plt.plot(train_sizes, val_scores_mean, 'o-', color='g', label='Cross-
validation score')
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color='r')
plt.fill_between(train_sizes, val_scores_mean - val_scores_std,
                 val_scores_mean + val_scores_std, alpha=0.1,
                 color='g')
plt.title('Learning Curves for HGBClassifier')
plt.xlabel('Training Size')
plt.ylabel('Accuracy')
plt.legend(loc='best')
plt.grid()
plt.show()

```



For discussion on overfitting, please see section 7.1 Technical Perspective

7. Final recommendation of the best model

7.1 Technical perspective- overfitting discussion, complexity and efficiency

7.1.1 Performance and Metrics

1. Overall Accuracy

- MLP: 0.8225 (about 82.25%)
- HGB: 0.8774 (about 87.74%)

From this, HGB is outperforming MLP in terms of simple accuracy.

2. Classification Report Details

MLP

- Class 0 - No Fraud:
 - Precision: 0.87
 - Recall: 0.89 (the model correctly identifies 89% of the "Class 0" cases)
 - F1-score: 0.88 - overall good performance
- Class 1 Fraud:
 - Precision: 0.69
 - Recall: 0.63 (the model correctly identifies 63% of the "Class 1" cases)
 - F1-score: 0.66 - overall score of precision and recall is lower

Overall, MLP tends to do better on Class 0, but struggles more with Class 1.

HGB

- Class 0:
 - Precision: 0.89
 - Recall: 0.94 (the model correctly identifies 94% of the "Class 0" cases)
 - F1-score: 0.92
- Class 1:
 - Precision: 0.82
 - Recall: 0.70 (the model correctly identifies 70% of the "Class 1" cases)
 - F1-score: 0.75

HGB shows stronger performance for both classes overall—particularly it does a better job on Class 1 compared to MLP.

3. Confusion Matrices

MLP

- 3761: True negatives (Class 0 correctly identified)

- 450: False positives (Class 1 was predicted when it was actually 0)
- *574: False negatives (Class 0 was predicted when it was actually 1)*
- 983: True positives (Class 1 correctly identified)

HGB

- 3978: True negatives (Class 0 correctly identified)
- 233: False positives (Class 1 was predicted when it was actually 0)
- *474: False negatives (Class 0 was predicted when it was actually 1)*
- 1083: True positives (Class 1 correctly identified)

```
#Confusion matrix values
cm_mlp = np.array([
    [3761, 450],
    [574, 983]
])

cm_hgb = np.array([
    [3978, 233],
    [474, 1083]
])

class_names = ['Class 0', 'Class 1']

# Create subplots
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 5))

# Plot for MLP
im1 = axes[0].imshow(cm_mlp, interpolation='nearest', cmap='Blues')
axes[0].set_title("Confusion Matrix (MLP)")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# Set up ticks & labels for MLP heatmap
axes[0].set_xticks(np.arange(len(class_names)))
axes[0].set_yticks(np.arange(len(class_names)))
axes[0].set_xticklabels(class_names)
axes[0].set_yticklabels(class_names)

# Annotate cells
for i in range(cm_mlp.shape[0]):
    for j in range(cm_mlp.shape[1]):
```

```

        axes[0].text(j, i, cm_mlp[i, j],
                    ha="center", va="center", color="black")

fig.colorbar(im1, ax=axes[0])

# Plot for HGB
im2 = axes[1].imshow(cm_hgb, interpolation='nearest', cmap='Greens')
axes[1].set_title("Confusion Matrix (HGB)")
axes[1].set_xlabel("Predicted")
axes[1].set_ylabel("Actual")

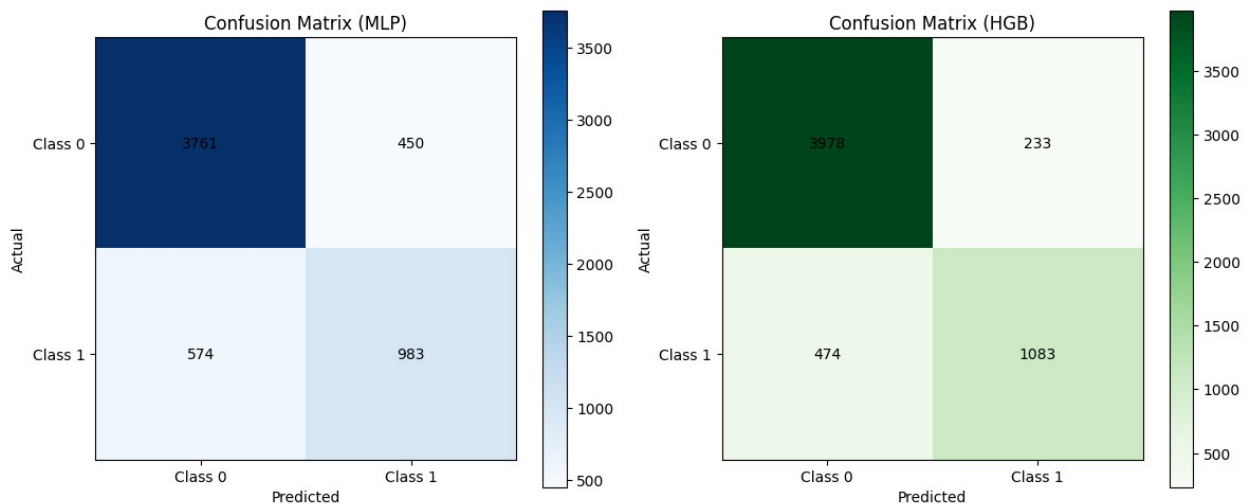
# Set up ticks & labels for HGB heatmap
axes[1].set_xticks(np.arange(len(class_names)))
axes[1].set_yticks(np.arange(len(class_names)))
axes[1].set_xticklabels(class_names)
axes[1].set_yticklabels(class_names)

# Annotate cells
for i in range(cm_hgb.shape[0]):
    for j in range(cm_hgb.shape[1]):
        axes[1].text(j, i, cm_hgb[i, j],
                    ha="center", va="center", color="black")

fig.colorbar(im2, ax=axes[1])

plt.tight_layout()
plt.show()

```



Comparing the matrices:

- HGB has fewer false positives (233 vs. 450) and fewer false negatives (474 vs. 574) than MLP.

- That means HGB is doing a better job of correctly labeling both Class 0 and Class 1 instances.
- In terms of fraud detection, the false negatives are more critical as they represent missed fraudulent claims. The false positives on the other hand are less critical as they represent genuine claims that were incorrectly flagged as fraudulent; these can be further investigated for verification.

7.1.2 Learning Curves & Overfitting

Learning Curves

HGB

- Training Score: Starts quite high (around 0.96) for small training sets and gradually decreases to the low 0.90s as training size grows.
- Cross-Validation (CV) Score: Starts around 0.82–0.83 and steadily increases to about 0.88.
- The training and CV curves move closer together as more data is used, suggesting the model is generalising well.
- There maybe some overfitting at small training sizes, where the training accuracy is much higher than the CV accuracy; but this gap closes as more data is used.
- Cross-Val (Training) Score: ~0.8756m Test Accuracy: 0.8774 - negligible gap between training and test accuracy, suggesting the model is not overfitting.
- At large training size, the F1 score is very similar between the test and training set, suggesting the model is not overfitting.

MLP

- Training Score: Fluctuates but is typically quite high (often above 0.90), peaking around 0.925 before dipping slightly toward the end. Cross-Validation (CV) Score: Starts around 0.75, climbs to the low 0.80s, and then slowly edges up to around 0.825.
- There is a larger and more persistent gap (about 0.10+ difference) between training and CV accuracy throughout most of the curve, although this does reduce with more data.
- There is a lot more fluctuation in the MLP curves, suggesting it may be more sensitive to the training data.
- The MLP model is more prone to overfitting, as indicated by the larger gap between training and CV scores.
- Cross-Val (Training) Score: ~0.8265 and Test Accuracy: 0.8225 That's only a 0.4% difference, which is very minor and suggests the model generalises well and not overfitting.
- At large training size, the F1 score is very similar between the test and training set, suggesting the model is not overfitting.

Precision Recall Curves

ROC Curves & AUC

- MLP: AUC = 0.84
- HGB: AUC = 0.89
- Area Under the ROC Curve (AUC) is a measure of how well the model separates the classes across different decision thresholds. HGB has a higher AUC, indicating better class separability.

Precision-Recall Curves & Average Precision (AP)

- MLP :AP = 0.81
- HGB :AP = 0.73
- Average Precision (AP) is a measure of the precision-recall trade-off and is representative of the area under the Precision-Recall curve. HGB has a higher AP, and stays above the blue curve (MLP) suggesting a stronger precision-recall balance.

7.1.3 Model Complexity and Efficiency

Complexity

- MLP: Being a neural network model, it can need more extensive tuning with learning rate, hidden layers, activation functions is more sensitive to hyperparameters and can be prone to overfitting.

HGB: Being a tree-based model, it is generally less sensitive to hyperparameters, more robust to outliers, and can handle imbalanced datasets effectively. It is also more memory-efficient and faster than GradientBoostingClassifier. Also much more straightforward to tune and less prone to overfitting.

Efficiency

- MLP: Training can be relatively slower for a large network or dataset, and it may require more computational resources. MLP can be regarded a black box model, making it harder to interpret the model. Inference can be fast once trained, but ultimately will depend on size.
- HGB: Training is faster and more memory-efficient, making it suitable for large datasets and real-time applications. Also easier to extract features importance and interpret the model, compared to MLP. Inference can be quicker and more efficient than MLP, as decision trees are simpler to evaluate.

7.1.4 Conclusion from Technical Perspective

- Both models generalise very well to the test set, showing very similar results to the cross-val score, with no sign of overfitting. This may have been in part due to early stopping and cross-validation.
- HGB not only has a higher overall accuracy (about 87.74% vs. 82.25%), but also better precision, recall, and F1-scores for both classes.
- The confusion matrix for HGB shows fewer mistakes in both labeling Class 0 and Class 1.

- HGB outperforms MLP in both ROC (AUC) and PR (AP) metrics, indicating better class separability and a stronger precision-recall trade-off.
- HGB's learning curves suggest it is more stable and generalises better than MLP, which shows more overfitting and sensitivity to the training data.
- HGB is more efficient and less complex than MLP, making it a better choice for large datasets and real-time applications, especially involving fraud detection using a binary classification model and where speed and accuracy are crucial. -HGB has a higher AUC and AP, indicating better class separability and a stronger precision-recall trade-off.
- Therefore, HGB is the stronger model in this comparison.

7.2 Business Perspective

The business perspective requirements:

- Find the average of all claims and calculate the gross profit for the company, assuming it is double the claims; this will require an annual profit calculation based on claims made for a fraction within the year
- Assuming 10% of customers base make claims, calculate the number of customers.
- Based on the pricing model, calculate the error cost of the models and the potential savings from the model.

7.2.1 Gross Profit Calculation

```
# calculate the total claim amount

property_claim = df_pivot['AmountOfPropertyClaim'].sum()
print(f'total property claim {property_claim}')

injury_claim = df_pivot['AmountOfInjuryClaim'].sum()
print(f'total injury claim {injury_claim}')

vehicle_claim = df_pivot['AmountOfVehicleDamage'].sum()
print(f'total vehicl claim {vehicle_claim}')

total_claim = property_claim + injury_claim + vehicle_claim
print(f'total claim {total_claim}')
```

```
total property claim 210037681
total injury claim 211573147
total vehicl claim 1086746063
total claim 1508356891
```

```
# find highest and lowest dateofincident
print(df_pivot['DateOfIncident'].min())
print(df_pivot['DateOfIncident'].max())
```

```
2015-01-01 00:00:00
2015-03-14 00:00:00
```

```
#calculate number of days between the highest and lowest date of incident
```

```
from datetime import datetime
date_format = "%Y-%m-%d"
a = datetime.strptime('2015-01-01', date_format)
b = datetime.strptime('2015-03-14', date_format)
delta = b - a
print(f'total number of days covered by data {delta.days}')
```

```
total number of days covered by data 72
```

```
#number of claims
```

```
print(f'number of claims {df_pivot['Unique_ID'].count()}')
```

```
number of claims 28836
```

```
#average number of incidents per day
```

```
print(f' average number of incidents per day  
{df_pivot['Unique_ID'].count()/delta.days}')
```

```
#average claim amount per day
```

```
print(f' average claim amount per day {total_claim/delta.days}')
```

```
average number of incidents per day 400.5
```

```
average claim amount per day 20949401.263888888
```

```
#estimated annual claim cost
```

```
estimated_annual_claim_cost = total_claim/delta.days * 365
print(f'estimated annual claim cost:  
{estimated_annual_claim_cost:.2f}')
```

```
company_profit_needed = estimated_annual_claim_cost * 2
```

```
print(company_profit_needed)
```

```
print(f"Company Profit_needed: {company_profit_needed:.2f}")
```

```
estimated annual claim cost: 7646531461.32
```

```
15293062922.638887
```

```
Company Profit_needed: 15293062922.64
```

Presumably this figure is in Indian rupees given all the incidents are in India.

7.2.2 Number of Customers

```
no_customers = df_pivot['Unique_ID'].count() * 10
print(f'no of customers: {no_customers.round()}')
```

```
# charge for each customer
```

```

premium_per_customer = company_profit_needed/no_customers
print(f'premium per customer: Indian rupees ₹
{charge_per_customer:.2f}')

```

```

no of customers: 288360
premium per customer: Indian rupees ₹ 53034.62

```

₹ 53034.62 amounts to £499.42 per customer, for an insurance policy which sound reasonable.

7.2.3 Error Cost Calculation

```

total_claim_amount = total_claim
num_claims = df_pivot['Unique_ID'].count()

# Model performance metrics (confusion matrix values)
fn_mlp = 574 # False negatives for MLP
fp_mlp = 450 # False positives for MLP

fn_gb = 474 # False negatives for Gradient Boosting
fp_gb = 233 # False positives for Gradient Boosting

# for annual error calculation
fn_mlp_annual = (fn_mlp / delta.days)*365
fp_mlp_annual = (fp_mlp / delta.days)*365

fn_gb_annual = (fn_gb / delta.days)*365
fp_gb_annual = (fp_gb / delta.days)*365

# Average Claim Amount per day
average_claim = total_claim_amount / num_claims #over 73 days
print("Average Claim Amount:", average_claim)

#Average Claim Amount per anum
average_claim_annum = average_claim * 365

# Step 2: Calculate Total Customers
total_customers = no_customers
print("Total Customers:", total_customers)

# Step 3: Calculate Premium Per Customer
print("Premium Per Customer:", premium_per_customer)

# Step 4: Calculate Losses for MLP
fn_loss_mlp = fn_mlp_annual * average_claim_annum # False negatives
fp_loss_mlp = fp_mlp_annual * 2 * average_claim_annum # False
positives
total_loss_mlp = fn_loss_mlp + fp_loss_mlp # Total loss

```

```

print("\nMLP Losses:")
print("FN Loss:", fn_loss_mlp)
print("FP Loss:", fp_loss_mlp)
print("Total Loss:", total_loss_mlp)

# Step 5: Calculate Losses for Gradient Boosting
fn_loss_gb = fn_gb_annual * average_claim_annum
fp_loss_gb = fp_gb_annual * 2 * average_claim_annum
total_loss_gb = fn_loss_gb + fp_loss_gb

print("\nGradient Boosting Losses:")
print("FN Loss:", fn_loss_gb)
print("FP Loss:", fp_loss_gb)
print("Total Loss:", total_loss_gb)

Average Claim Amount: 52308.11801220696
Total Customers: 288360
Premium Per Customer: 53034.61965126539

MLP Losses:
FN Loss: 55556415815.68307
FP Loss: 87109362777.2034
Total Loss: 142665778592.88647

Gradient Boosting Losses:
FN Loss: 45877597729.327126
FP Loss: 45103292282.418655
Total Loss: 90980890011.74579

```

In summary:

Overall Loss Comparison:

- MLP: ₹ 142.66B total loss
- HGB: ₹ 90.98B total loss

MLP

- FN Loss: 55.56B
- FP Loss: 87.11B

HGB

- FN Loss: 45.88B
- FP Loss: 45.10B

HGB losses are balanced whereas MLP has a higher loss from false positives, suggesting it is overestimating the number of fraudulent claims. Most importantly, the FN, which is the most critical error in fraud detection, is lower in HGB, suggesting it is better at identifying fraudulent claims.

HGB is significantly more cost-effective than MLP, with a potential saving of ₹ 51.68B; due to HGB's better performance in correctly identifying fraudulent claims, leading to fewer false negatives and more savings for the company.

7.3.4 Potential Savings and Loss from the Models

#If HGB chosen over MLP

```
Savings_from_HGB = total_loss_mlp - total_loss_gb  
print(f"Savings from HGB over MLP: ₹ {Savings_from_HGB:.2f}")
```

Savings from HGB over MLP: ₹ 51684888581.14

How much HGB can cost the company due to false negatives

```
loss_from_HGB = fn_loss_gb  
print(f"Loss from HGB due to False Negatives: ₹ {loss_from_HGB:.2f}")
```

Loss from HGB due to False Negatives: ₹ 45877597729.33

#Model performance metrics (confusion matrix values)

```
tp_mlp = 983 # True positive for MLP  
tn_mlp = 3761 # True negative for MLP
```

```
tn_gb = 3978 # True negatives for Gradient Boosting  
tp_gb = 983 # True positives for Gradient Boosting
```

for annual error calculation

```
tn_mlp_annual = (tn_mlp / delta.days)*365  
tp_mlp_annual = (tp_mlp / delta.days)*365
```

```
tn_gb_annual = (tn_gb / delta.days)*365  
tp_gb_annual = (tp_gb / delta.days)*365
```

#Calculate Gains for Gradient Boosting

```
tn_gain_gb = tn_gb_annual * average_claim_annum # True negatives  
tp_gain_gb = tp_gb_annual * 2 * average_claim_annum # True positives  
total_gain_gb = tn_gain_gb + tp_gain_gb # Total gain
```

```
print("\nGradient Boosting Positives:")  
print("TN Positive:", tn_gain_gb)  
print("TP Positive:", tp_gain_gb)  
print("Total Positives:", total_gain_gb)
```

```
Gradient Boosting Positives:  
TN Positive: 385023383475.2391  
TP Positive: 190285563577.7577  
Total Positives: 575308947052.9968
```

```
#How much a company can save by using HGB over no model

#by using HGB, the company can correctly identify fraudulent causes

savings_due_to_HGB = tp_gain_gb #identify fraud causes
print(f"Savings due to recognising fraud claims using HGB: ₹
{savings_due_to_HGB:.2f}")

total_savings = savings_due_to_HGB - loss_from_HGB #total savings
print(f"Total Savings using HGB: ₹ {total_savings:.2f}")

Savings due to recognising fraud claims using HGB: ₹ 190285563577.76
Total Savings using HGB: ₹ 144407965848.43
```

7.3.5 Conclusion from Business Perspective

- The company will save ₹ 144.4BN by using HGB; this includes loss due to false negatives and gains due to true positives.
- HGB is more cost-effective than MLP, with a potential saving of ₹ 51.68B over MLP due to HGB's better performance in correctly identifying fraudulent claims, leading to fewer false negatives and more savings for the company.

7.3 Overall Recommendation

HistGradientBoostingClassifier is recommended over MLPClassifier for the following reasons:

- HGB has a higher overall accuracy, precision, recall, and F1-scores for both classes.
- HGB shows fewer mistakes in both labeling Class 0 and Class 1.
- HGB outperforms MLP in both ROC (AUC) and PR (AP) metrics, indicating better class separability and a stronger precision-recall trade-off.
- HGB's learning curves suggest it is more stable and generalises better than MLP, which shows more overfitting and sensitivity to the training data.
- HGB is more efficient and less complex than MLP, making it a better choice for large datasets and real-time applications, especially involving fraud detection using a binary classification model and where speed and accuracy are crucial.
- HGB is significantly more cost-effective than MLP, with a potential saving of ₹ 51.68B over MLP due to HGB's better performance in correctly identifying fraudulent claims, leading to fewer false negatives and more savings for the company.

Therefore, HGB is the stronger model in this comparison and is recommended for fraud detection in the insurance company.

8. Conclusion

8.1 What has been successfully accomplished and what has not

After weeks of non-stop work, I have successfully pre-processed the data, built two models, and evaluated them using nested cross-validation. I have also compared the models from both technical and business perspectives. The HistGradientBoostingClassifier model outperformed

the MLPClassifier in terms of accuracy, precision, recall, F1-scores, and cost-effectiveness. The HGB model is more stable, generalises better, and is more efficient than the MLP model. It is also more cost-effective, with a potential saving of ₹ 51.68B over MLP due to its better performance in correctly identifying fraudulent claims.

The aims and objectives in the project brief have been met, and the final recommendation is to use the HGB model for fraud detection in the insurance company.

What hasn't been accomplished is the unbiased predictive model having a balanced error of 5% or less.

8.2 Reflect back on the analysis and see what you could have done differently if you were to do the project again

- Balancing the weights was tricky using MLP, as the library didn't have built-in handling for class imbalance; this took significant time to resolve. In hindsight, I could have used a different model that has built-in handling for class imbalance.
- I could have used a different feature selection method, such as PCA, to see if it would have improved the model's performance.
- I also underestimated the time it would take to resolve technical issues throughout the model building and nested cross-validation process. I could have allocated more time for this. I had to use RandomizedSearchCV instead of GridSearchCV for the MLP model due to the time constraints as GridSearchCV was taking too long to run.

8.3 Add a wish list of future work that you would do to take the project forward

As future work, I would be interested in exploring the following:

- How Bayesian optimization could be used to tune the hyperparameters of the models more efficiently; also how this compares to GridSearchCV and RandomizedSearchCV.
- If there is scope for unsupervised learning techniques, such as anomaly detection, to identify fraudulent claims without the need for labels.
- How ensemble methods, such as stacking, could be used to combine the predictions of multiple models to improve the overall performance.
- Find models that have a lower error rate, approximating the 5% error rate target set in the project brief.

9. References

- Ahmad, I. (2023) 50 Algorithms Every Programmer Should Know: Tackle Computer Science Challenges with Classic to Modern Algorithms in Machine Learning, Software Design, Data Systems, and Cryptography. 2nd edn. Birmingham: Packt Publishing.
- Géron, A. (2022) Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems. 3rd edn. Sebastopol, CA: O'Reilly Media.

- Raschka, S., Liu, Y. (H.) and Mirjalili, V. (2022) Machine Learning with PyTorch and Scikit-Learn: Develop Machine Learning and Deep Learning Models with Python. Birmingham: Packt Publishing.
- Tan, P.-N., Steinbach, M., Karpatne, A. and Kumar, V. (2018) Introduction to Data Mining. 2nd edn. Boston: Pearson Addison Wesley.

Websites

- [Scikit-learn: Model Evaluation](#)
- [Scikit-learn: Grid Search](#)
- [Scikit-learn: MLPClassifier](#)
- [Scikit-learn: HistGradientBoostingClassifier](#)
- [Scikit-learn: Feature Selection](#)
- [Imbalanced-learn: SMOTE](#)
- [Imbalanced-learn: ClusterCentroids](#)
- [Imbalanced-learn: Pipeline](#)
- [Imbalanced-learn: RandomUnderSampler](#)
- [Imbalanced-learn: RandomOverSampler](#)

Acknowledgement

- Generative AI - ChatGPT- 01 by OpenAI (2021) [OpenAI](#)
- Context: Support tutor to aid understanding, debug code, and trouble shoot errors; research internet on resources needed for the project.